While you're waiting...

http://bit.ly/srecon-cake0

Kubernetes From A Cake Mix

—

Liz Frost

Who am I?

Liz Frost

Twitter: @stillinbeta

slack.k8s.io: @liz

Who am I?

Who am I?

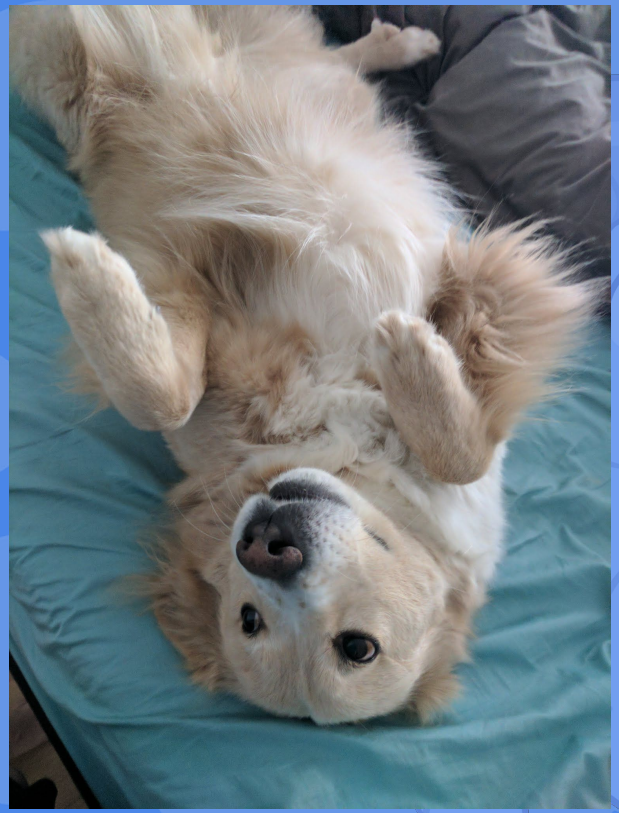

Who am I?

It’s pronounced
“kube-cuddle”

Why should you listen to me?

- sig-cluster-lifecycle member
- kubernetes contributor
- kubeadm developer

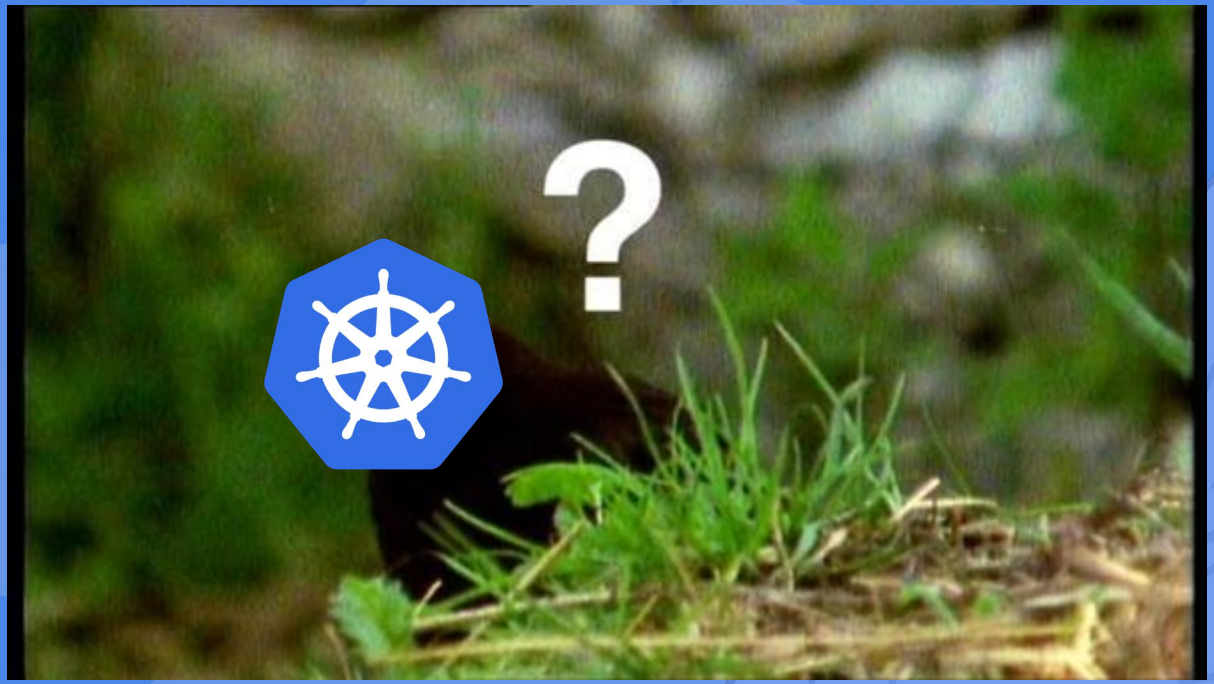

What *is* Kubernetes?


Node

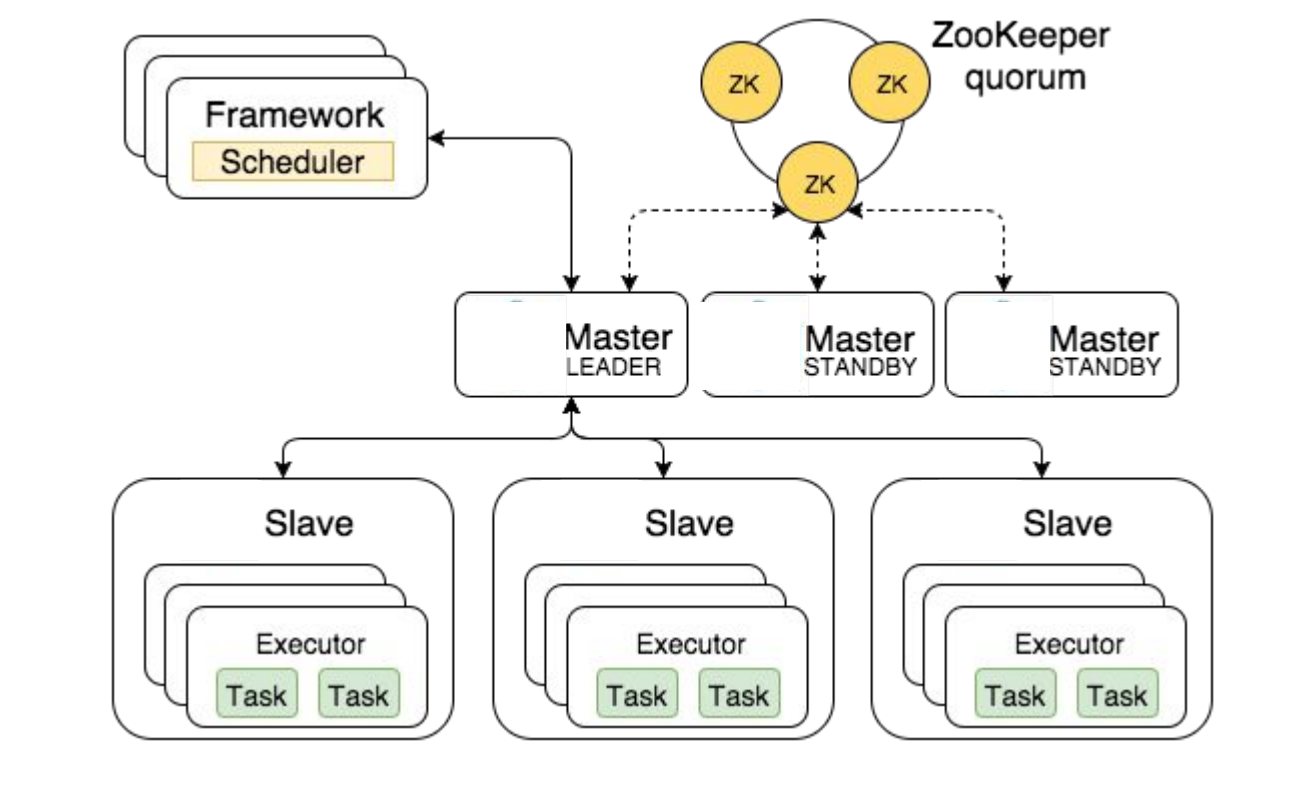
ZooKeeper
quorum
ZK
ZK
ZK
Framework
Scheduler
Master
LEADER
Master
STANDBY
Master
STANDBY
Slave
Executor
Task
Task
Slave
Executor
Task
Task
Slave
Executor
Task
Task

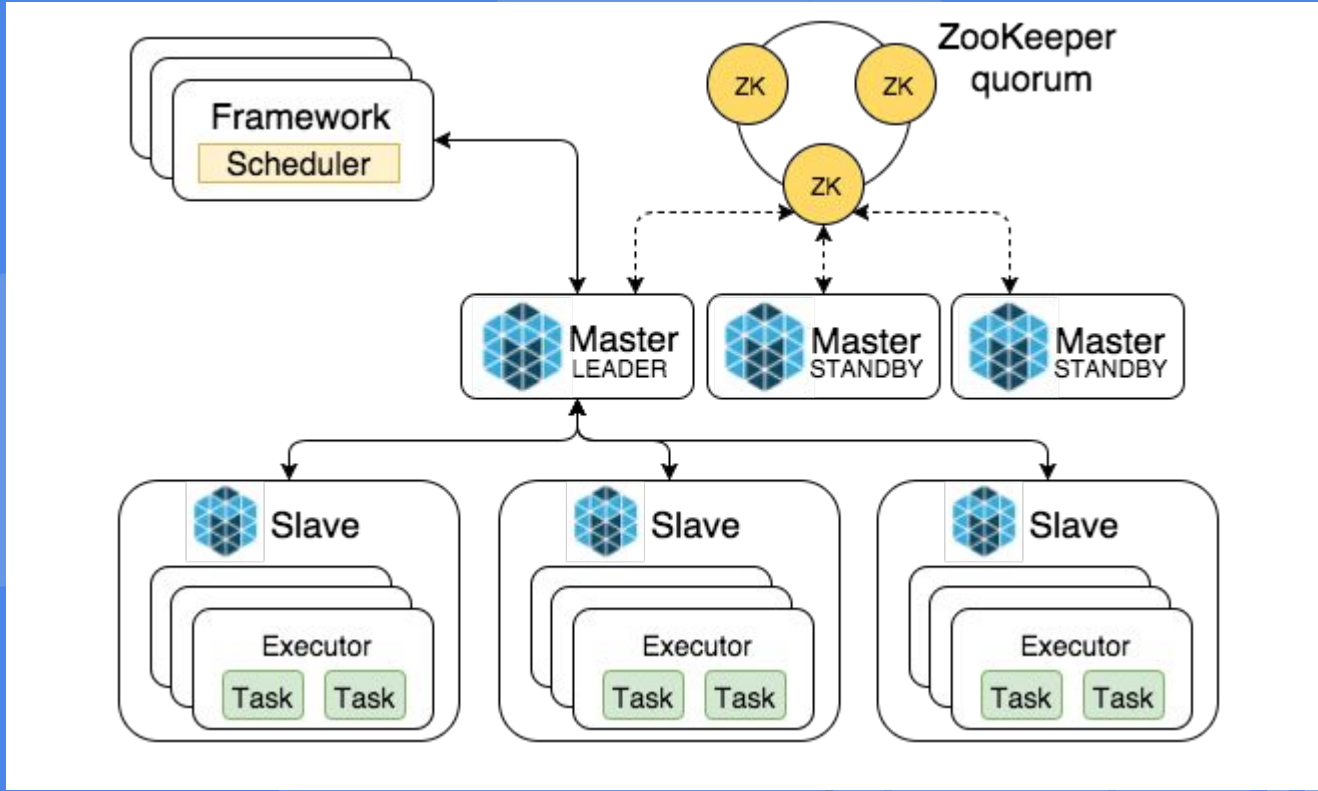
Framework
Scheduler
ZK
ZK
ZK
ZooKeeper
quorum
Master
LEADER
Master
STANDBY
Master
STANDBY
Slave
Executor
Task
Task
Slave
Executor
Task
Task
Slave
Executor
Task
Task

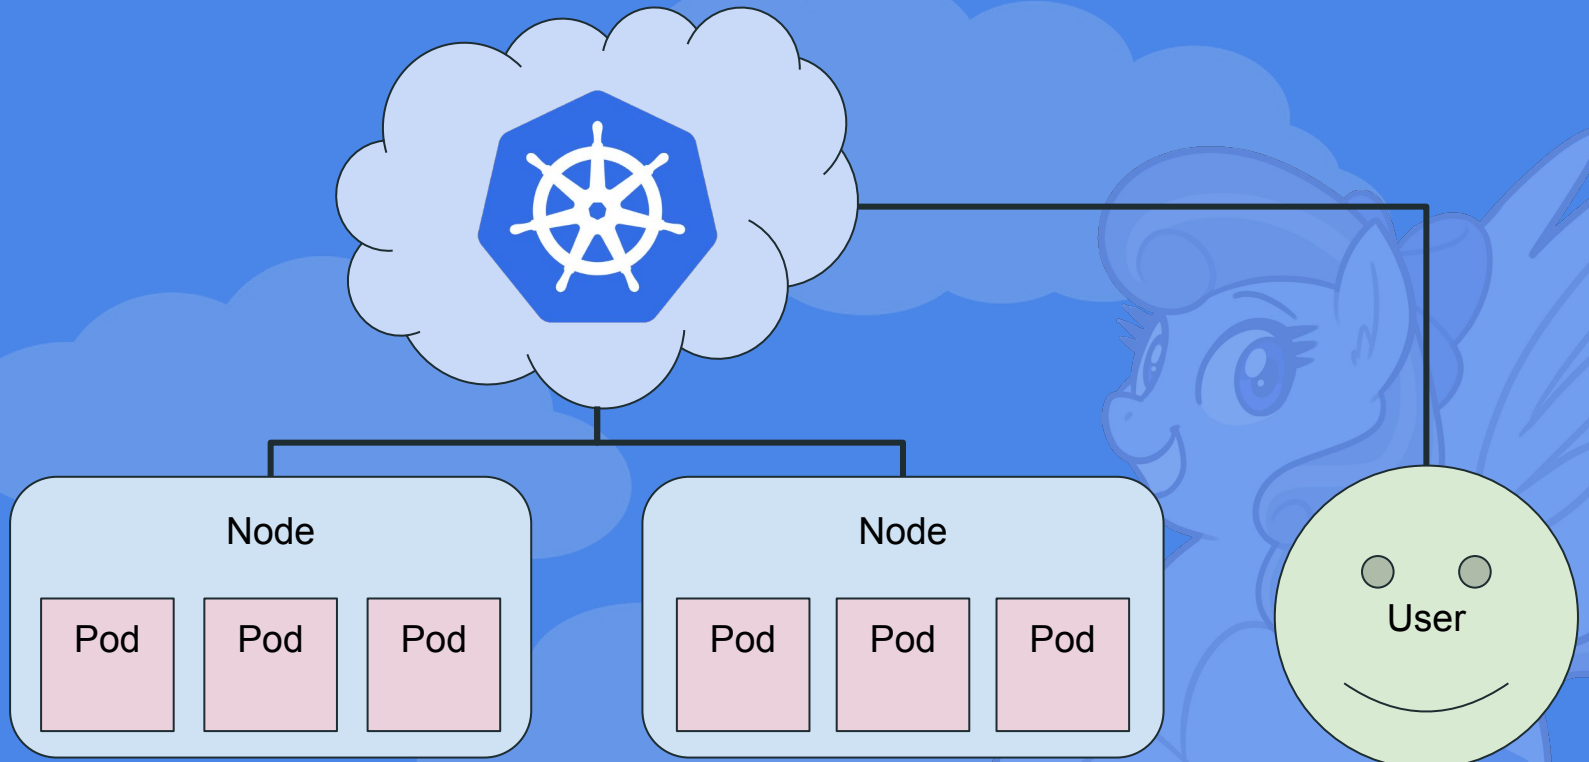
Node
Pod
Pod
Pod
Node
Pod
Pod
Pod
User

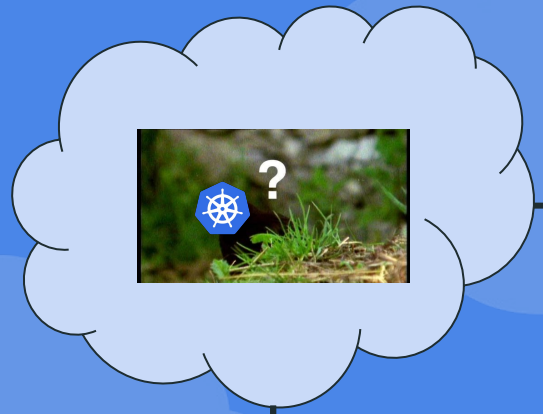

Node

Pod

Pod

Pod

Node

Pod

Pod

Pod

User

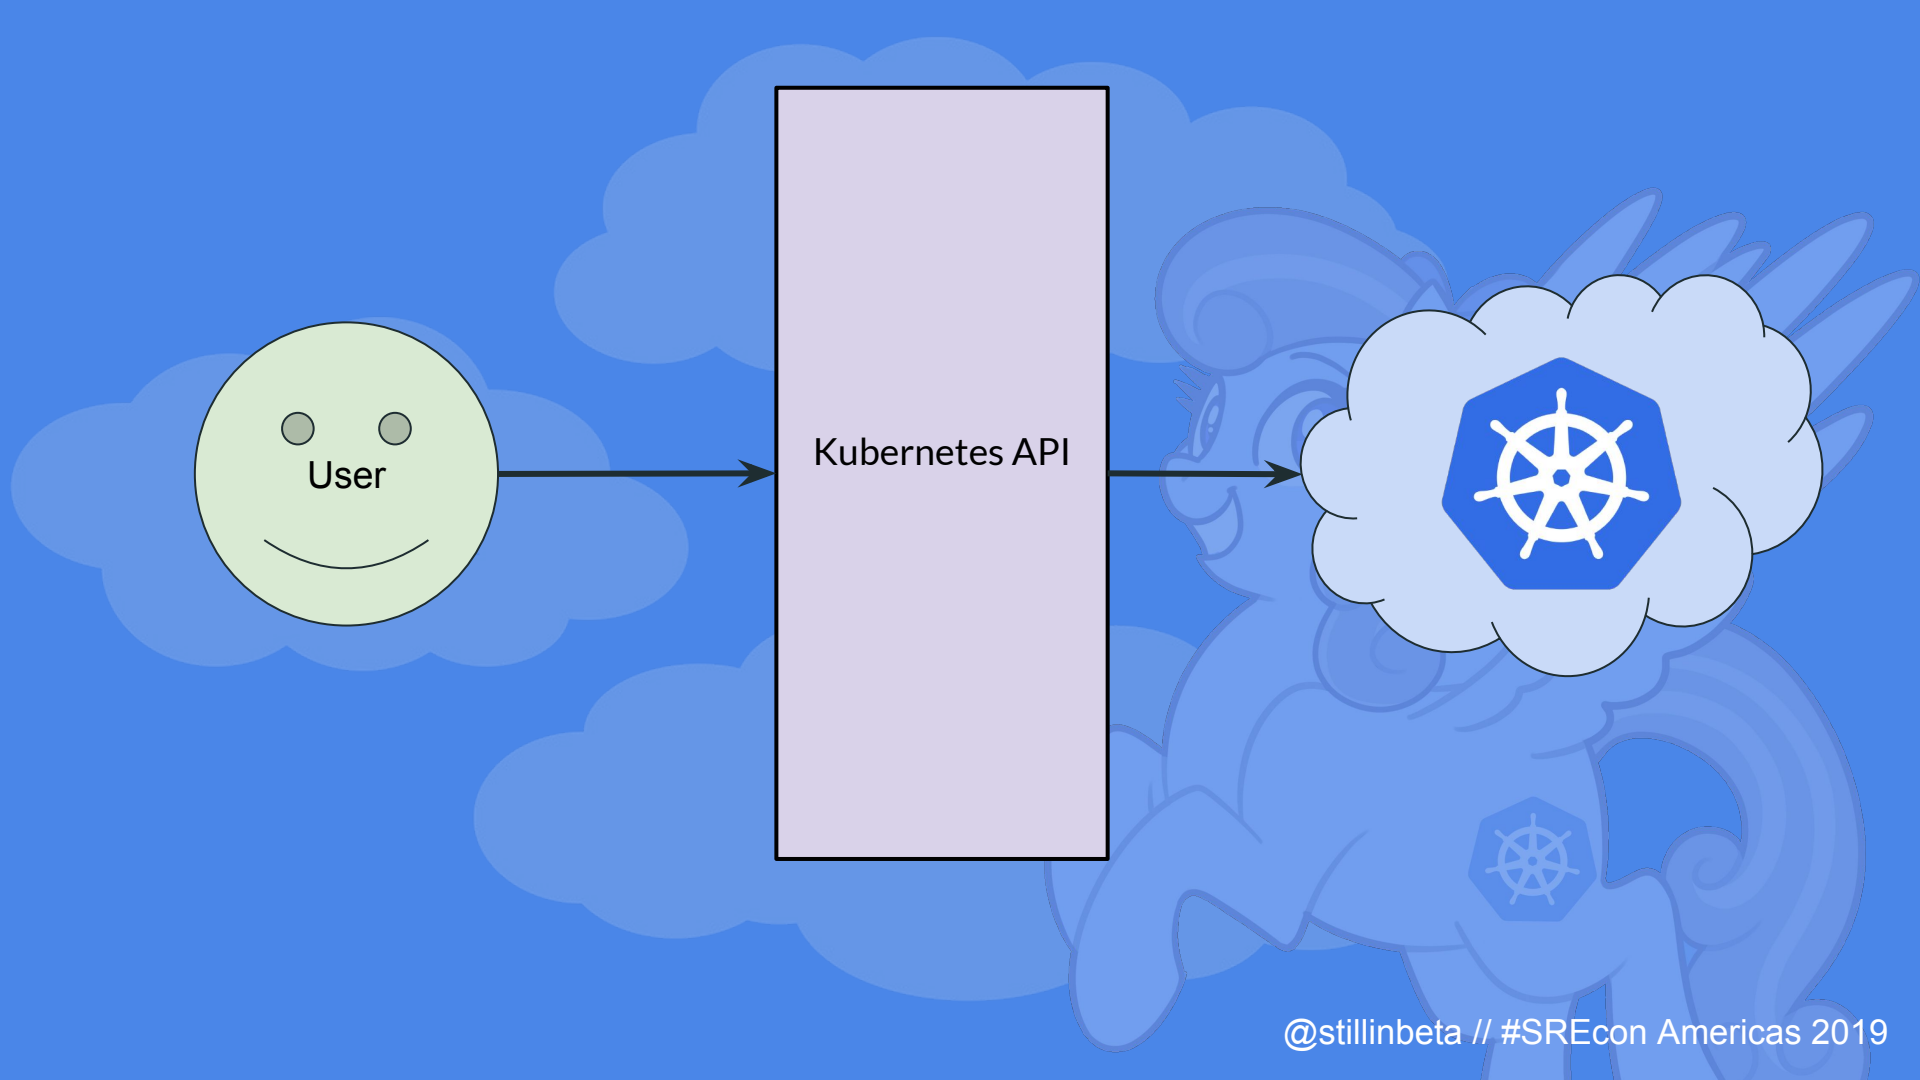
User
Kubernetes API

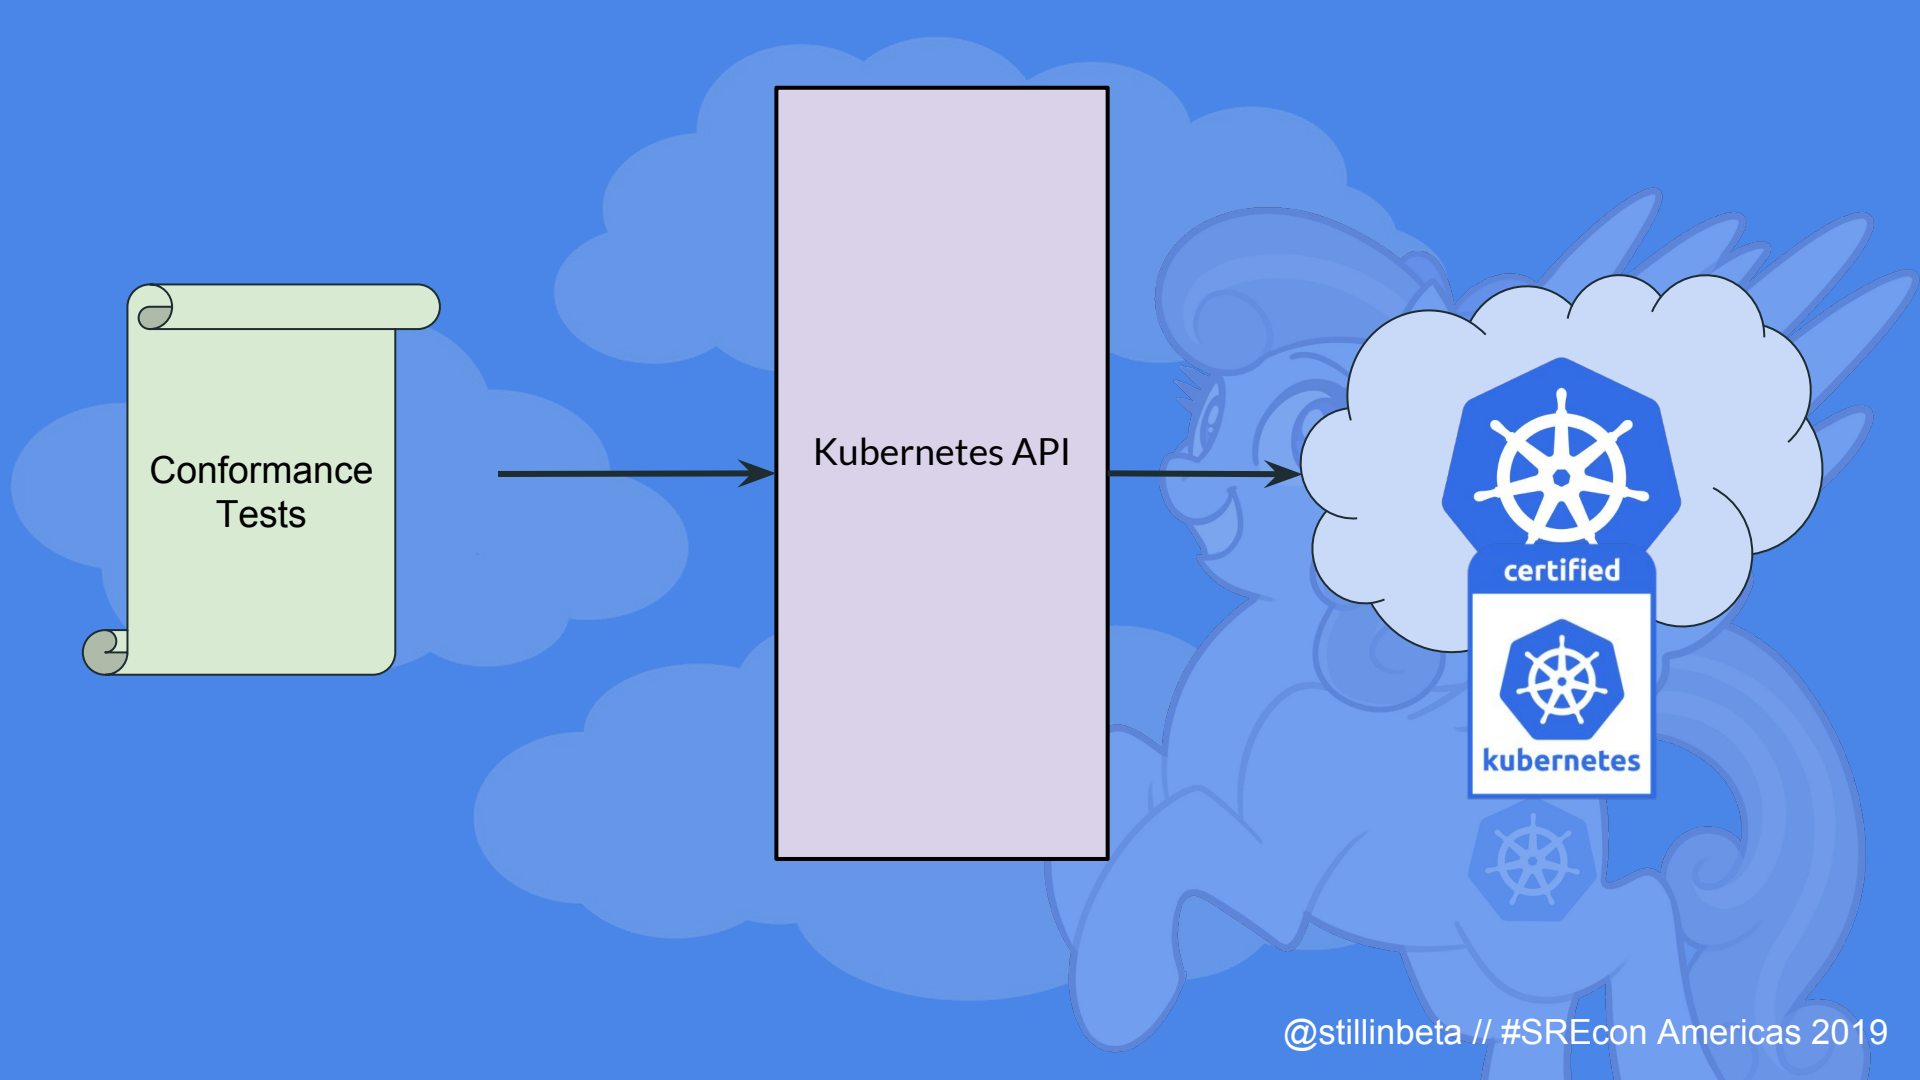
Conformance Tests
Kubernetes API
certified
kubernetes

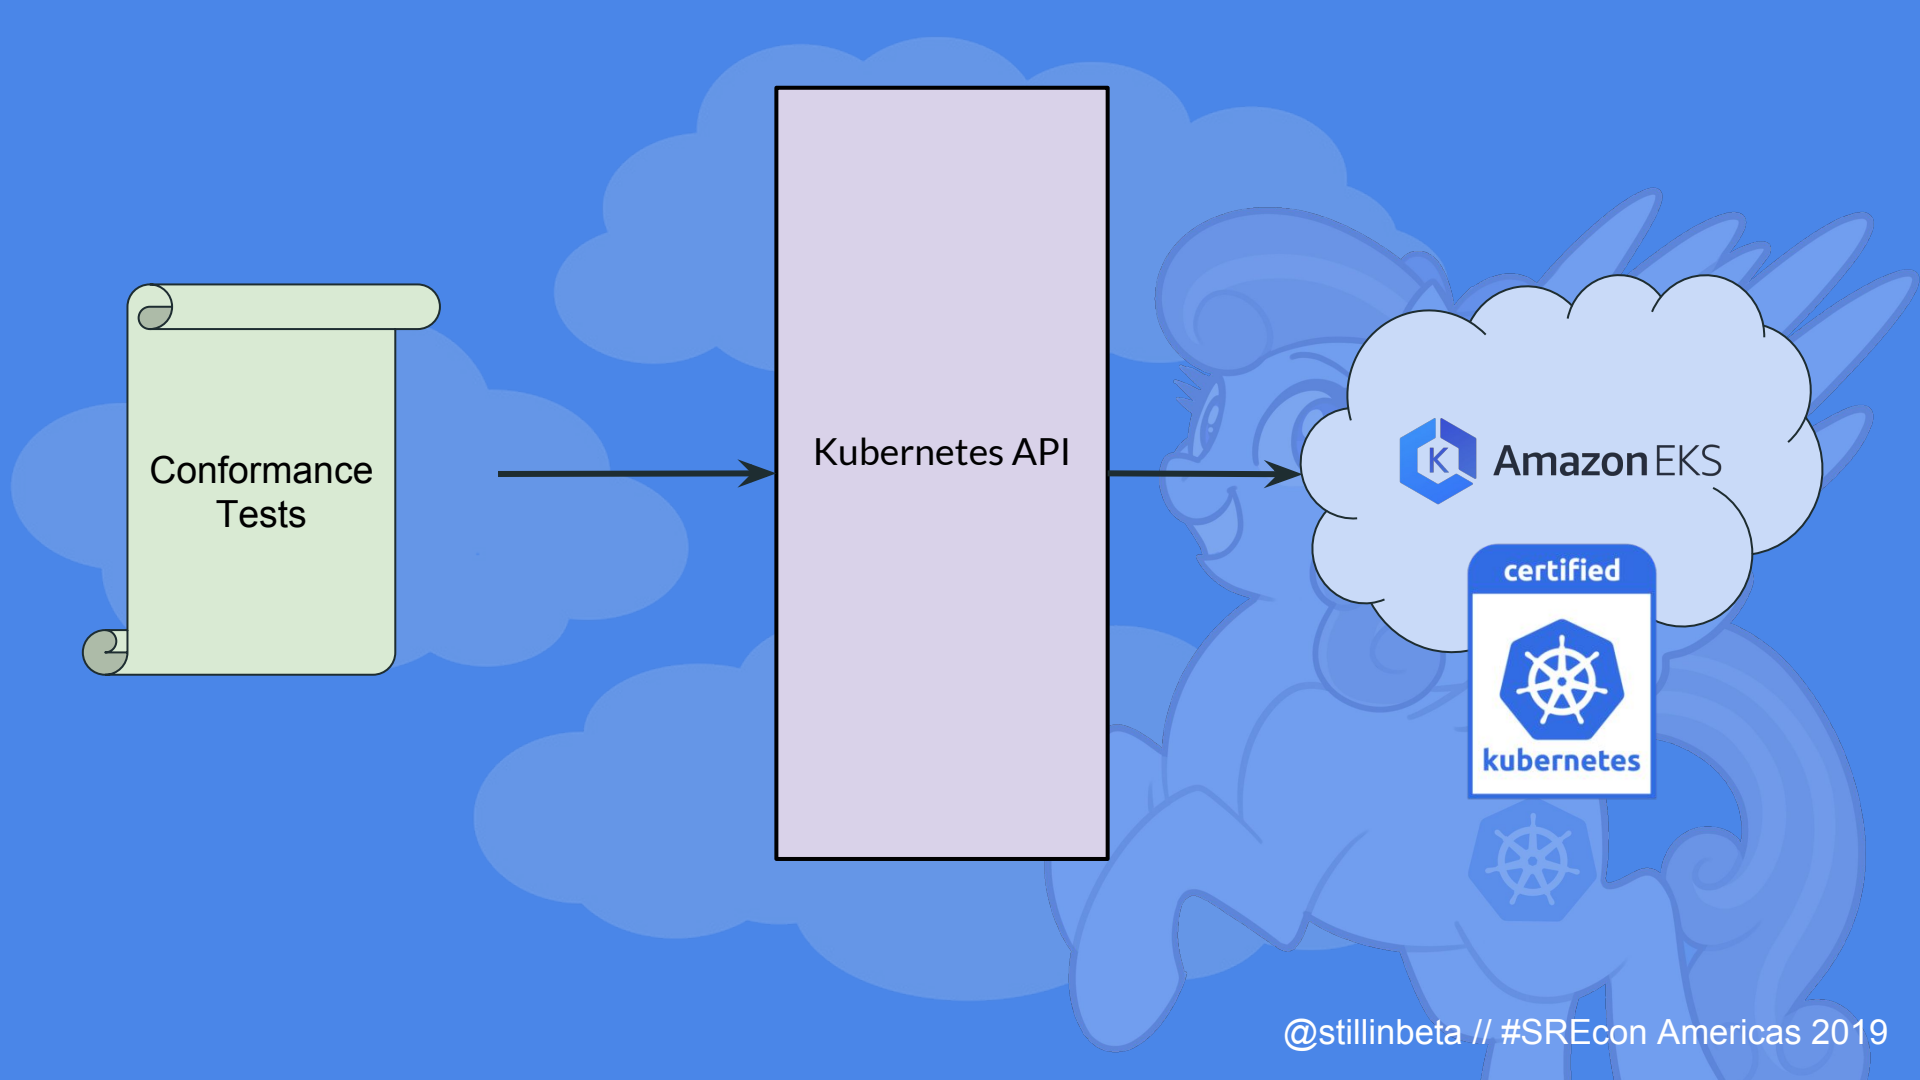
Conformance Tests
Kubernetes API
Amazon EKS
certified
kubernetes

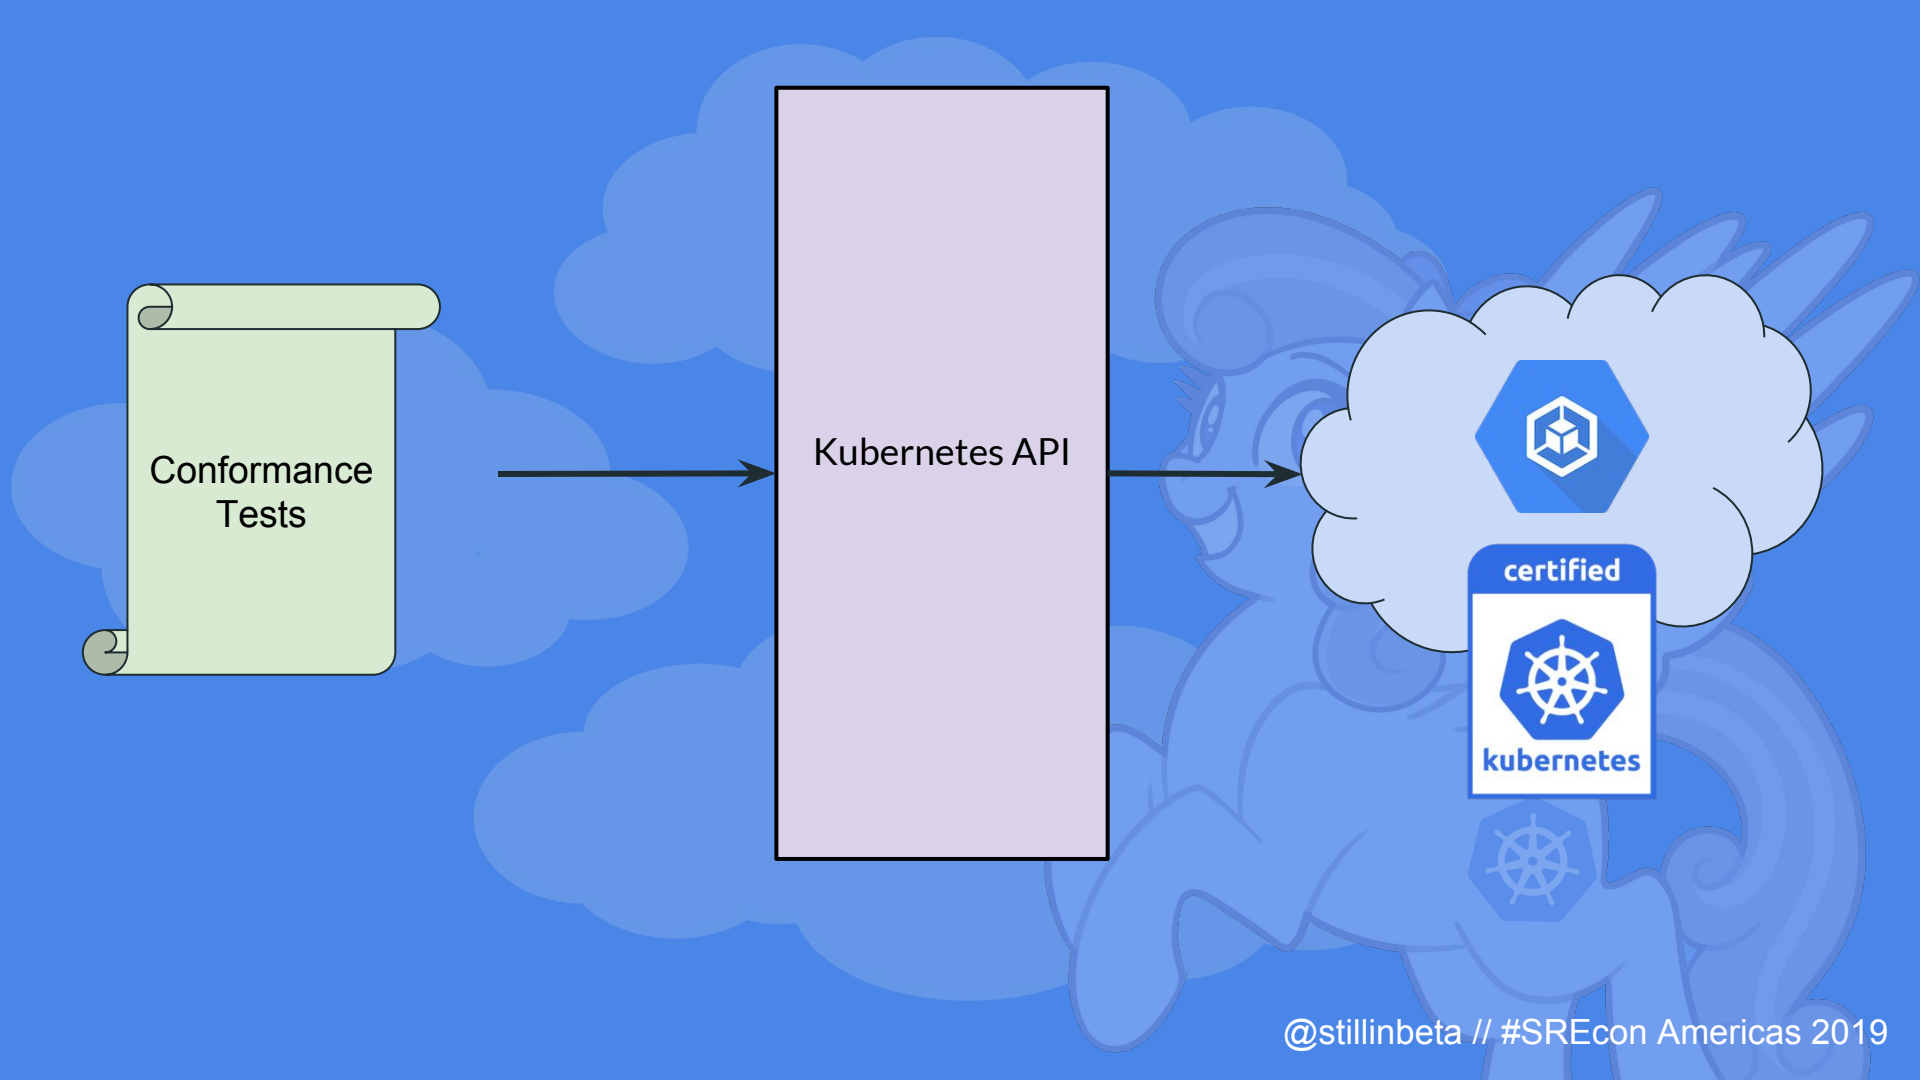
Conformance Tests
Kubernetes API
certified
kubernetes

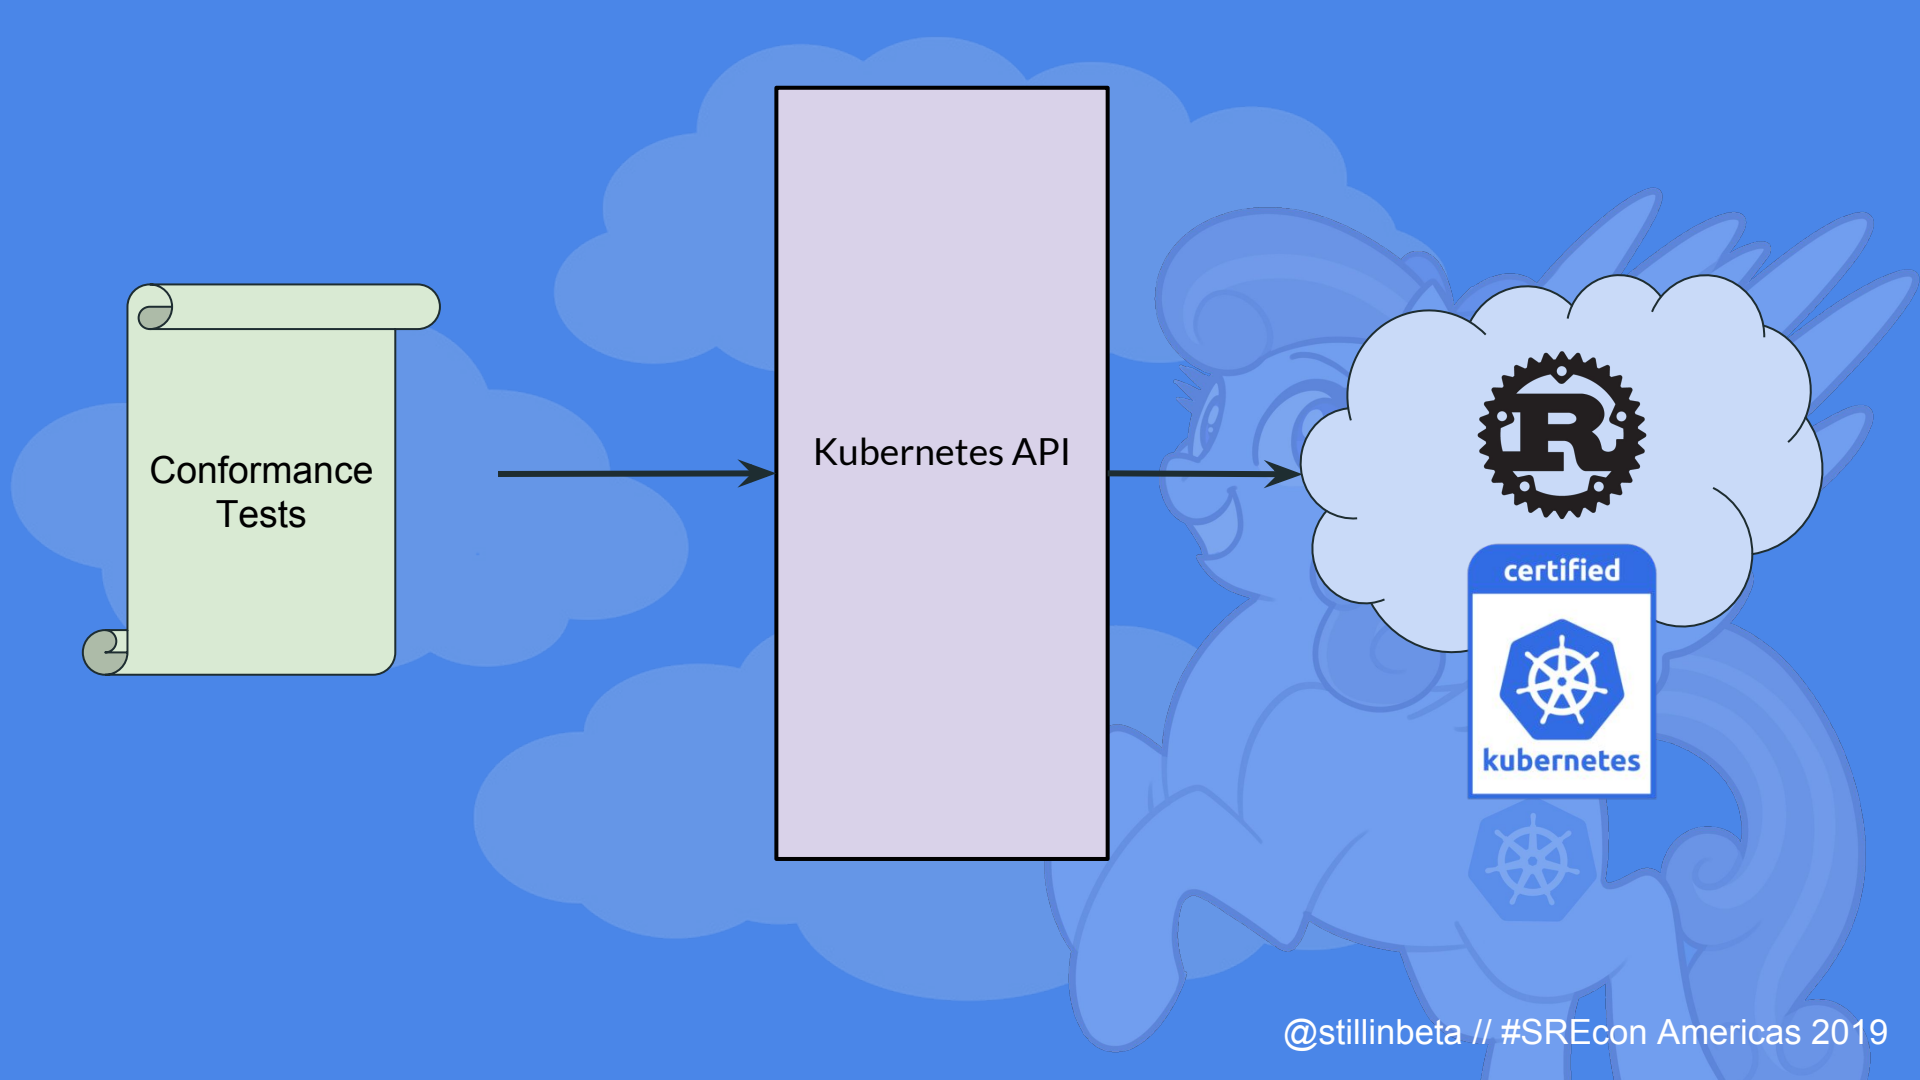
Conformance Tests
Kubernetes API
certified
kubernetes

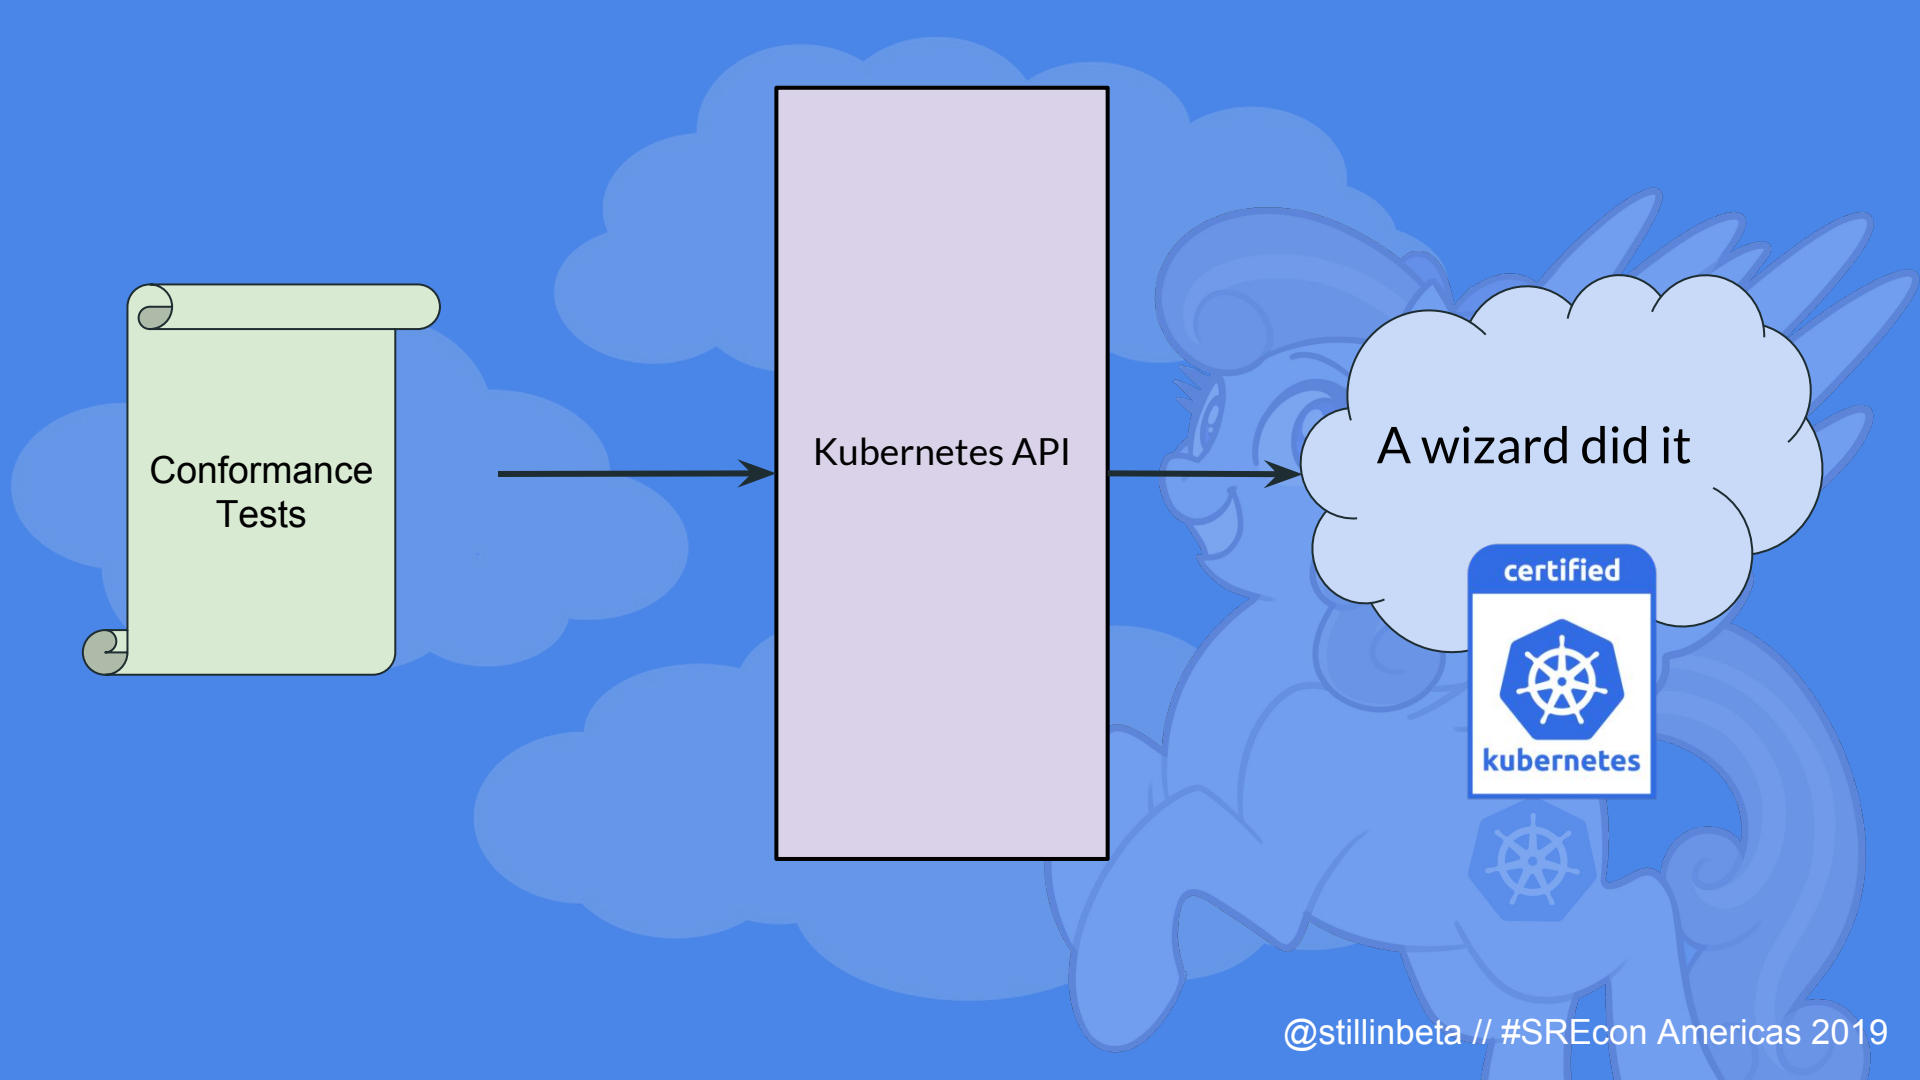
Conformance Tests
Kubernetes API
A wizard did it
certified
kubernetes

We got Opinions


NEW!
SIGNATURE
Kubeadm™ brand
Kubernetes

kubeadm init

```
$ kubeadm init --help
Run this command in order to set up the Kubernetes master.

The "init" command executes the following phases:
```

```
preflight                  Run master pre-flight checks
kubelet-start              Writes kubelet settings and (re)starts the kubelet
certs                      Certificate generation
  /etcd-ca                   Generates the self-signed CA to provision identities for etcd
  /etcd-server               Generates the certificate for serving etcd
  /etcd-healthcheck-client   Generates the client certificate for liveness probes to healtcheck etcd
  /apiserver-etcd-client     Generates the client apiserver uses to access etcd
  /etcd-peer                 Generates the credentials for etcd nodes to communicate with each other
  /ca                        Generates the self-signed Kubernetes CA to provision identities for other Kubernetes components
  /apiserver                 Generates the certificate for serving the Kubernetes API
  /apiserver-kubelet-client  Generates the Client certificate for the API server to connect to kubelet
  /front-proxy-ca            Generates the self-signed CA to provision identities for front proxy
  /front-proxy-client        Generates the client for the front proxy
  /sa                        Generates a private key for signing service account tokens along with its public key
kubeconfig                 Generates all kubeconfig files necessary to establish the control plane and the admin kubeconfig
file
  /admin                     Generates a kubeconfig file for the admin to use and for kubeadm itself
  /kubelet                   Generates a kubeconfig file for the kubelet to use *only* for cluster bootstrapping purposes
  /controller-manager        Generates a kubeconfig file for the controller manager to use
  /scheduler                 Generates a kubeconfig file for the scheduler to use
control-plane              Generates all static Pod manifest files necessary to establish the control plane
  /apiserver                 Generates the kube-apiserver static Pod manifest
  /controller-manager        Generates the kube-controller-manager static Pod manifest
  /scheduler                 Generates the kube-scheduler static Pod manifest
etcd                       Generates static Pod manifest file for local etcd.
  /local                     Generates the static Pod manifest file for a local, single-node local etcd instance.
upload-config              Uploads the kubeadm and kubelet configuration to a ConfigMap
  /kubeadm                   Uploads the kubeadm ClusterConfiguration to a ConfigMap
  /kubelet                   Uploads the kubelet component config to a ConfigMap
mark-control-plane         Mark a node as a control-plane
bootstrap-token            Generates bootstrap tokens used to join a node to a cluster
addon                      Installs required addons for passing Conformance tests
  /coredns                   Installs the CoreDNS addon to a Kubernetes cluster
  /kube-proxy                Installs the kube-proxy addon to a Kubernetes cluster
```

```
$ kubeadm init --help
Run this command in order to set up the Kubernetes master.

The "init" command executes the following phases:
```

```
preflight                   Run master pre-flight checks
kubelet-start               Writes kubelet settings and (re)starts the kubelet
certs                       Certificate generation
kubeconfig                  Generates all kubeconfig files necessary
control-plane               Generates all static Pod manifest files necessary to establish the
  /apiserver                  Generates the kube-apiserver static Pod manifest
  /controller-manager         Generates the kube-controller-manager static Pod manifest
  /scheduler                  Generates the kube-scheduler static Pod manifest
etcd                        Generates static Pod manifest file for local etcd.
  /local                      Generates the static Pod manifest file for a local etcd instance.
upload-config               Uploads the kubeadm and kubelet configuration to a ConfigMap
  /kubeadm                    Uploads the kubeadm ClusterConfiguration to a ConfigMap
  /kubelet                    Uploads the kubelet component config to a ConfigMap
mark-control-plane          Mark a node as a control-plane
bootstrap-token             Generates bootstrap tokens used to join a node to a cluster
addon                       Installs required addons for passing Conformance tests
  /coredns                    Installs the CoreDNS addon to a Kubernetes cluster
  /kube-proxy                 Installs the kube-proxy addon to a Kubernetes cluster
```

```
preflight
kubelet-start
certs
kubeconfig
control-plane
  /apiserver
  /controller-manager
  /scheduler
etcd
  /local
upload-config
  /kubeadm
  /kubelet
mark-control-plane
bootstrap-token
addon
  /coredns
  /kube-proxy
```

What are we making?

- `kubelet`
- `etcd`
- `kube-apiserver`
- `kube-contoller-manager`
- `kube-scheduler`

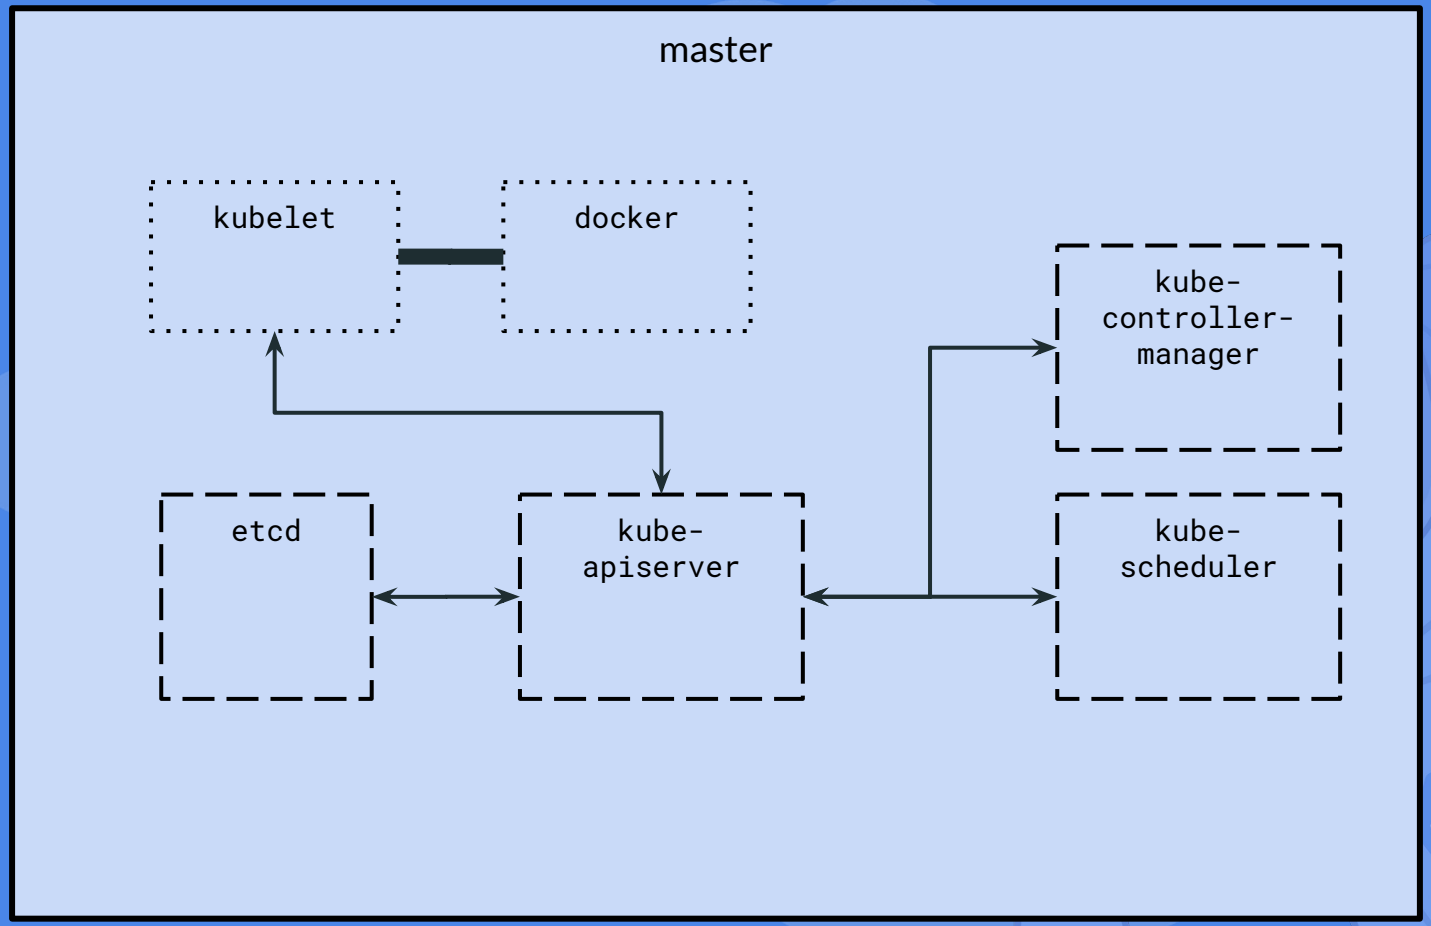

Node (VM / physical Machine)

process

container

https

IPC

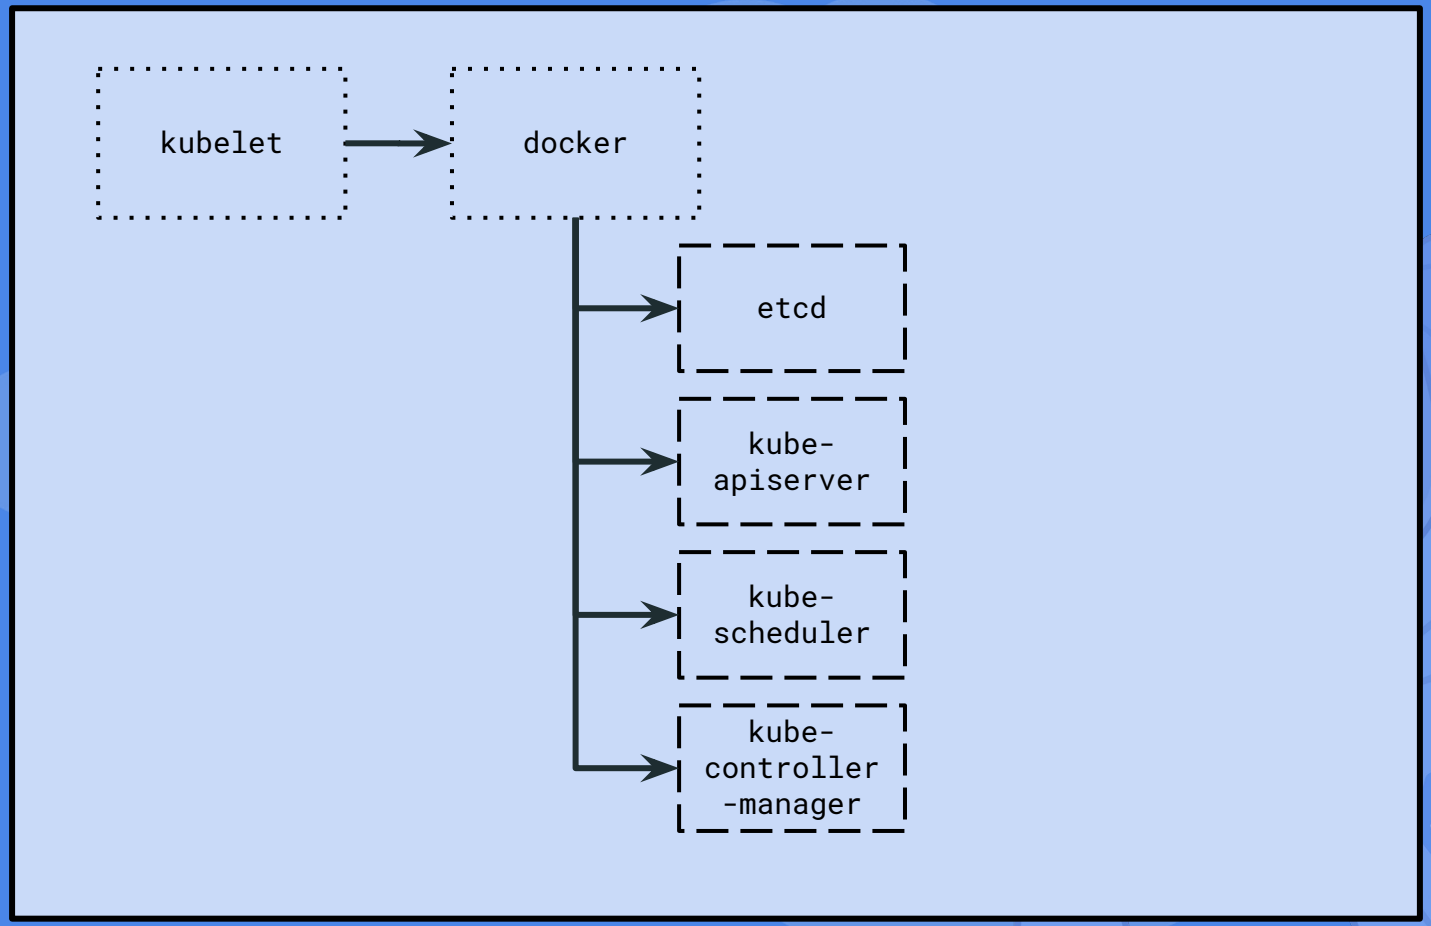

Node (VM / physical Machine)

process

container

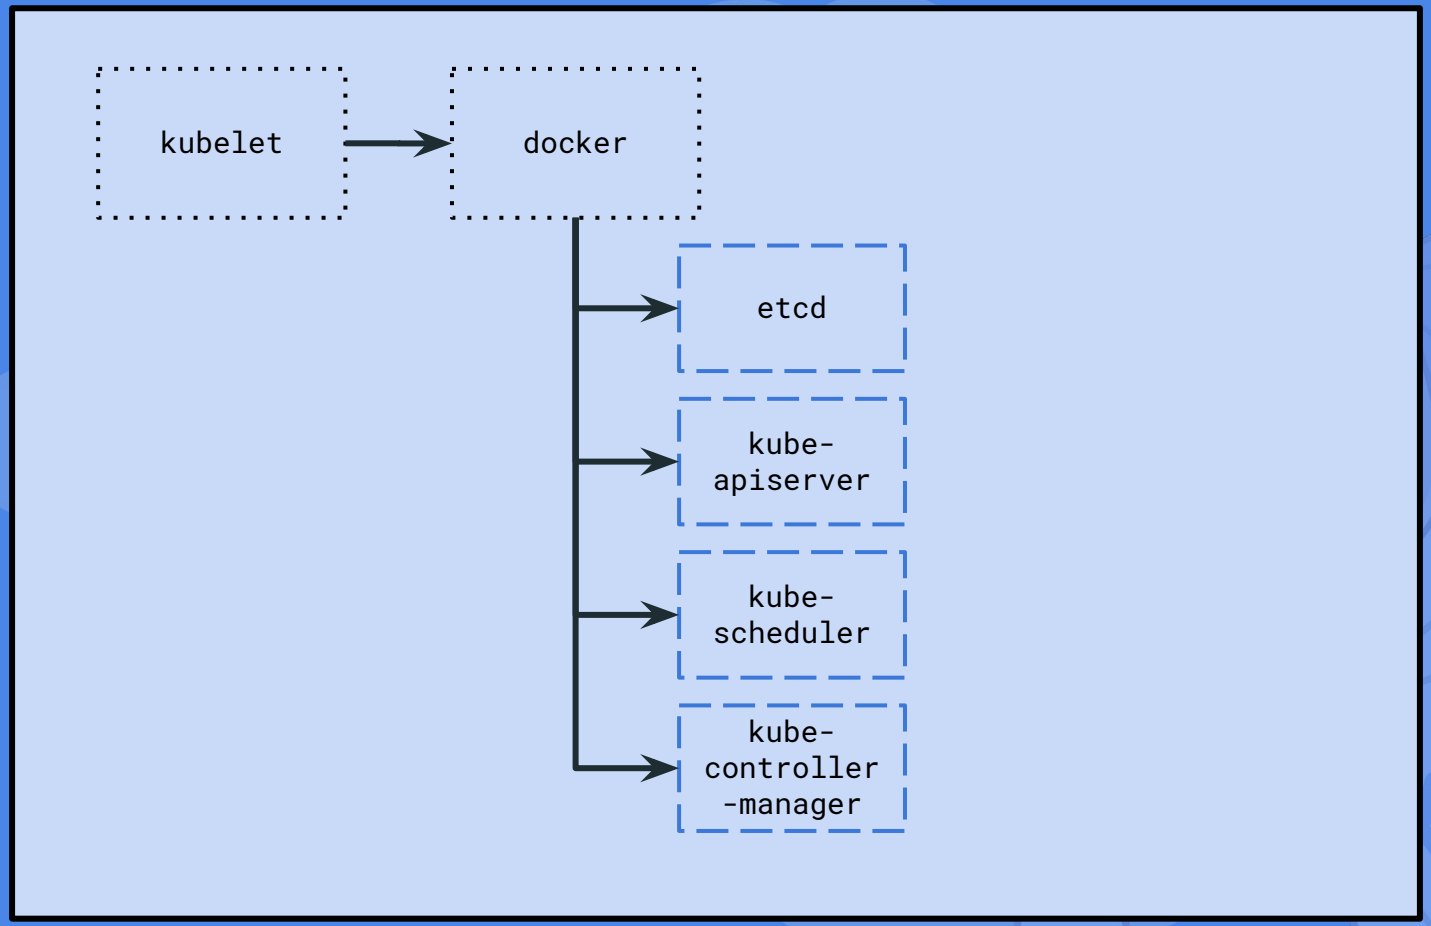

Node (VM / physical Machine)

process

Kubernetes pod

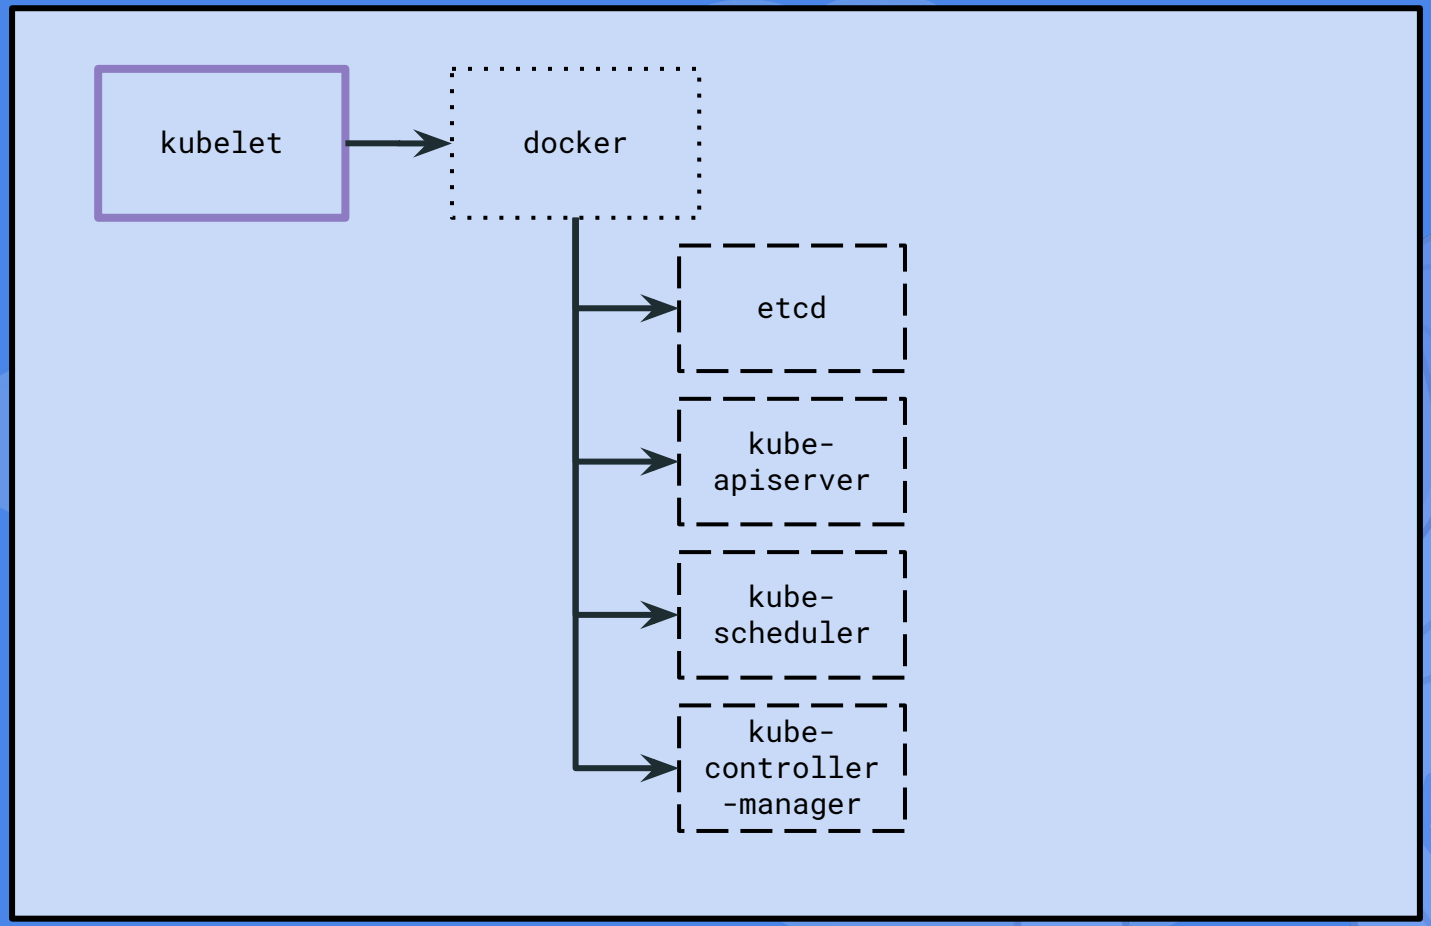

Node (VM / physical Machine)

process

container

kubelet

What will we need locally?

- kubelet
- kubeadm
- docker

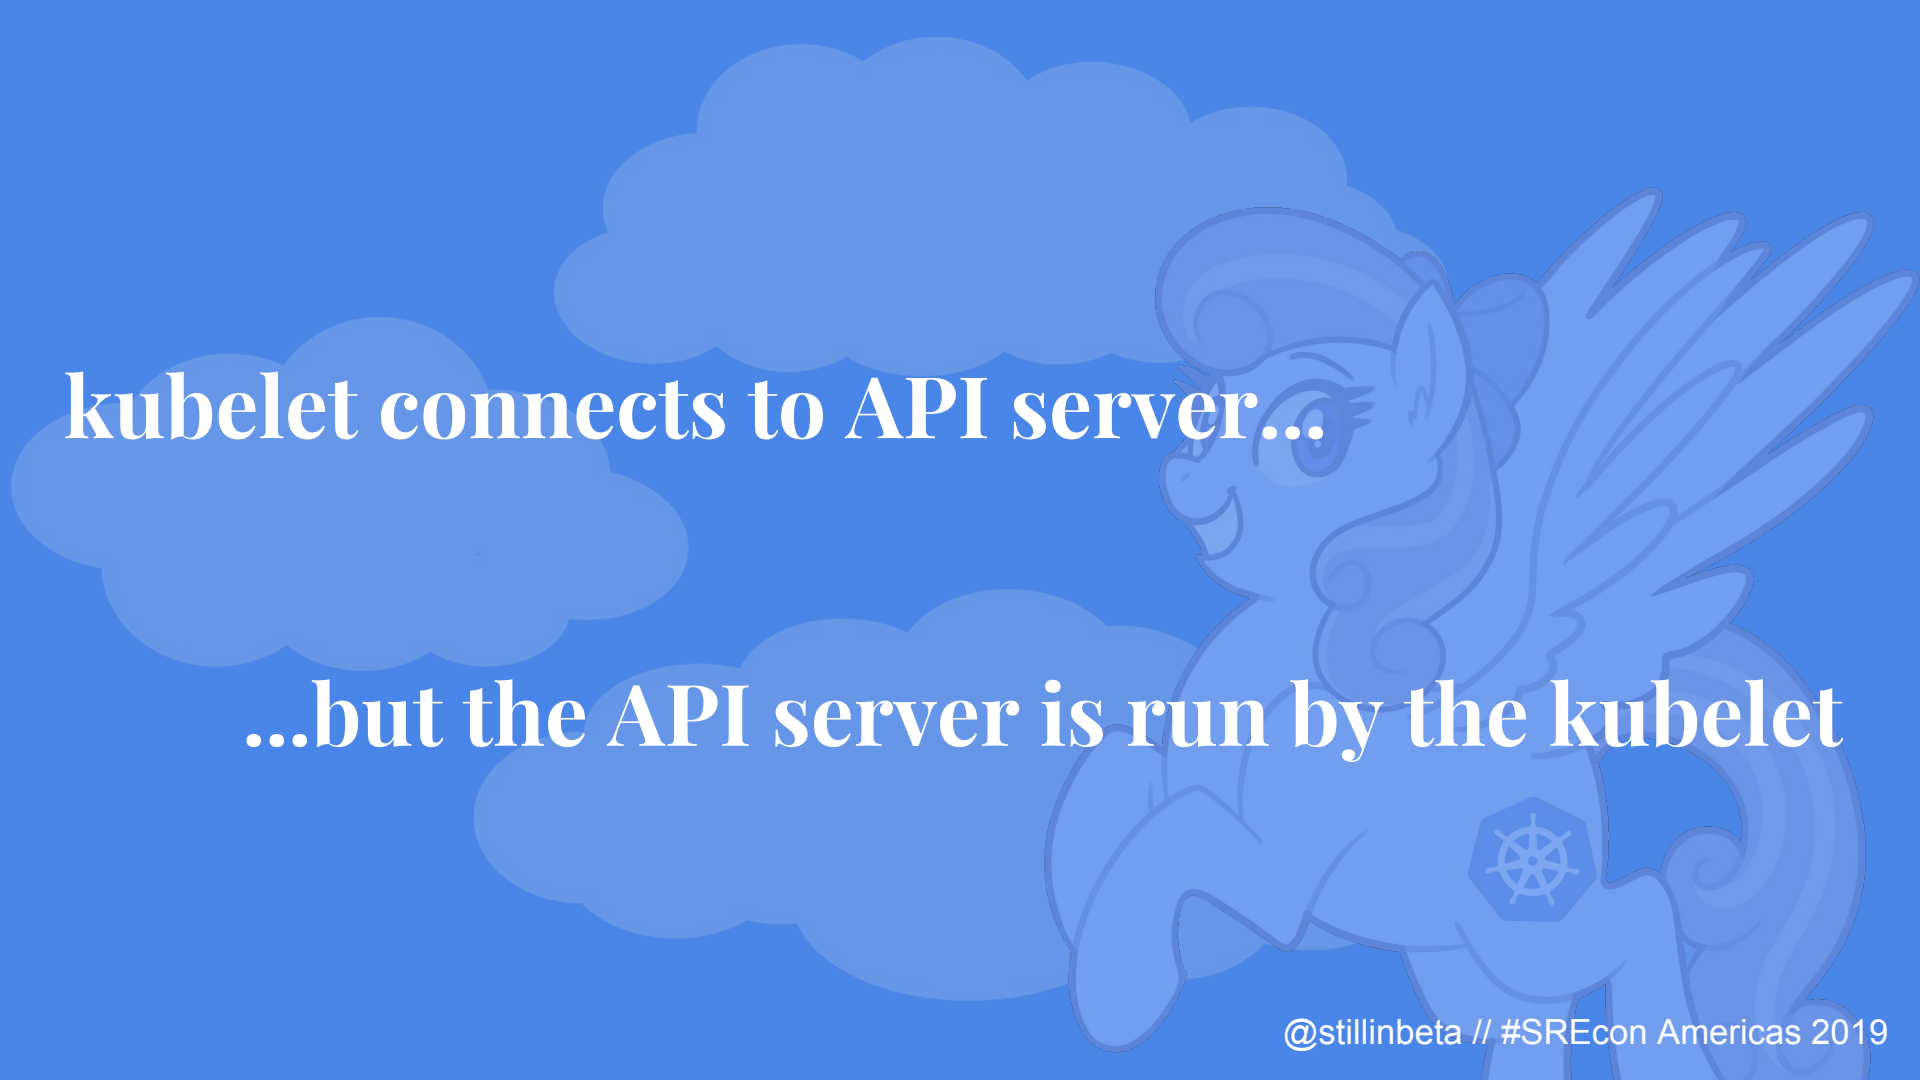
kubelet connects to API server...
...but the API server is run by the kubelet

MANIFEST

Kubelet boot

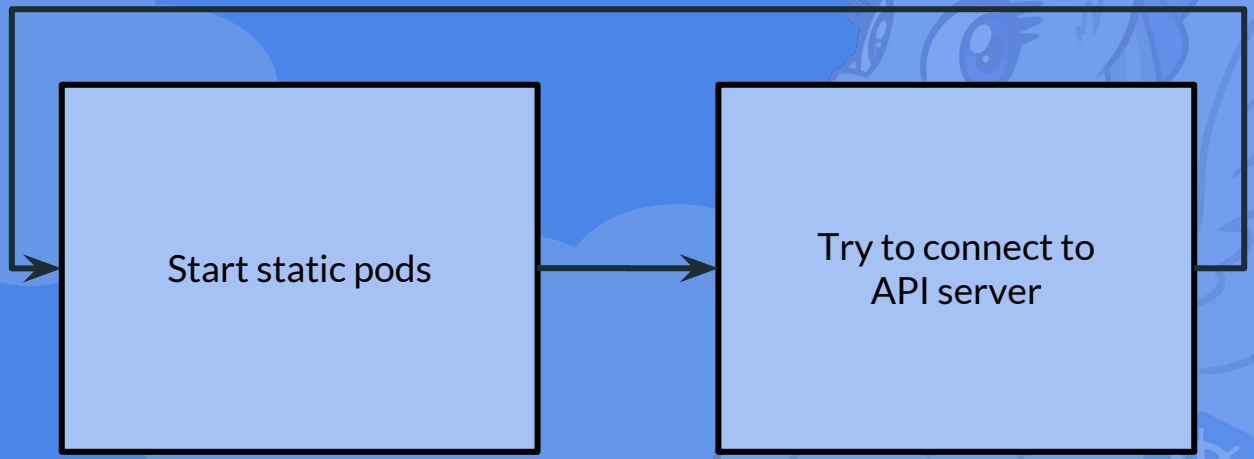

Kubelet boot

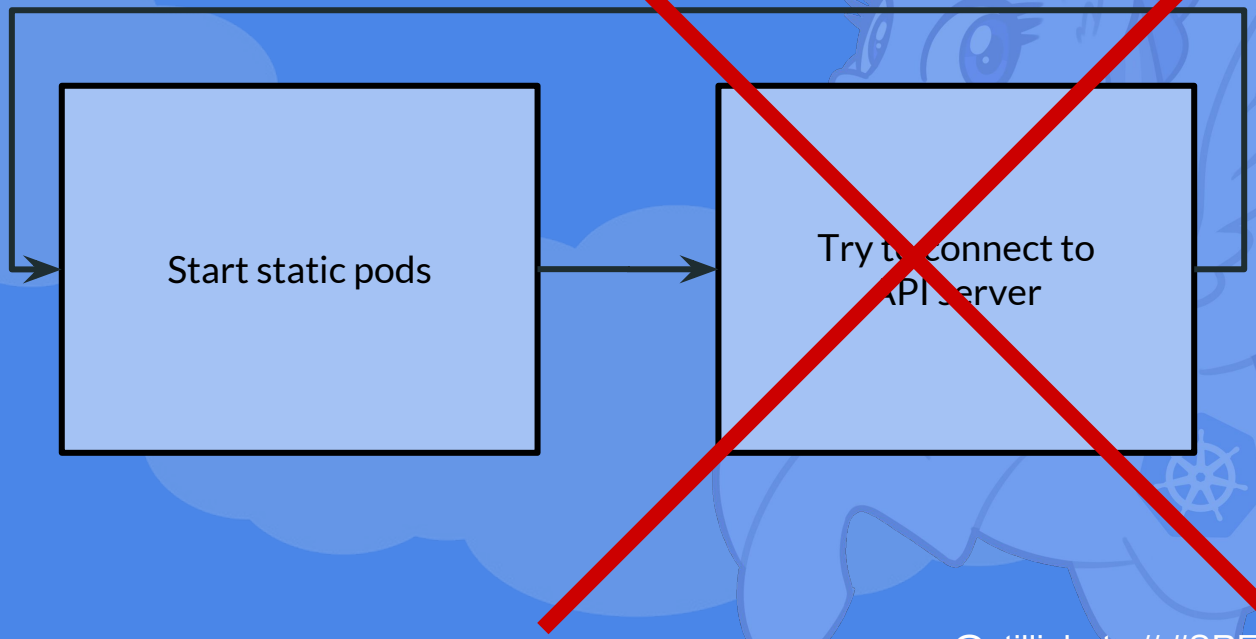

Kubelet boot

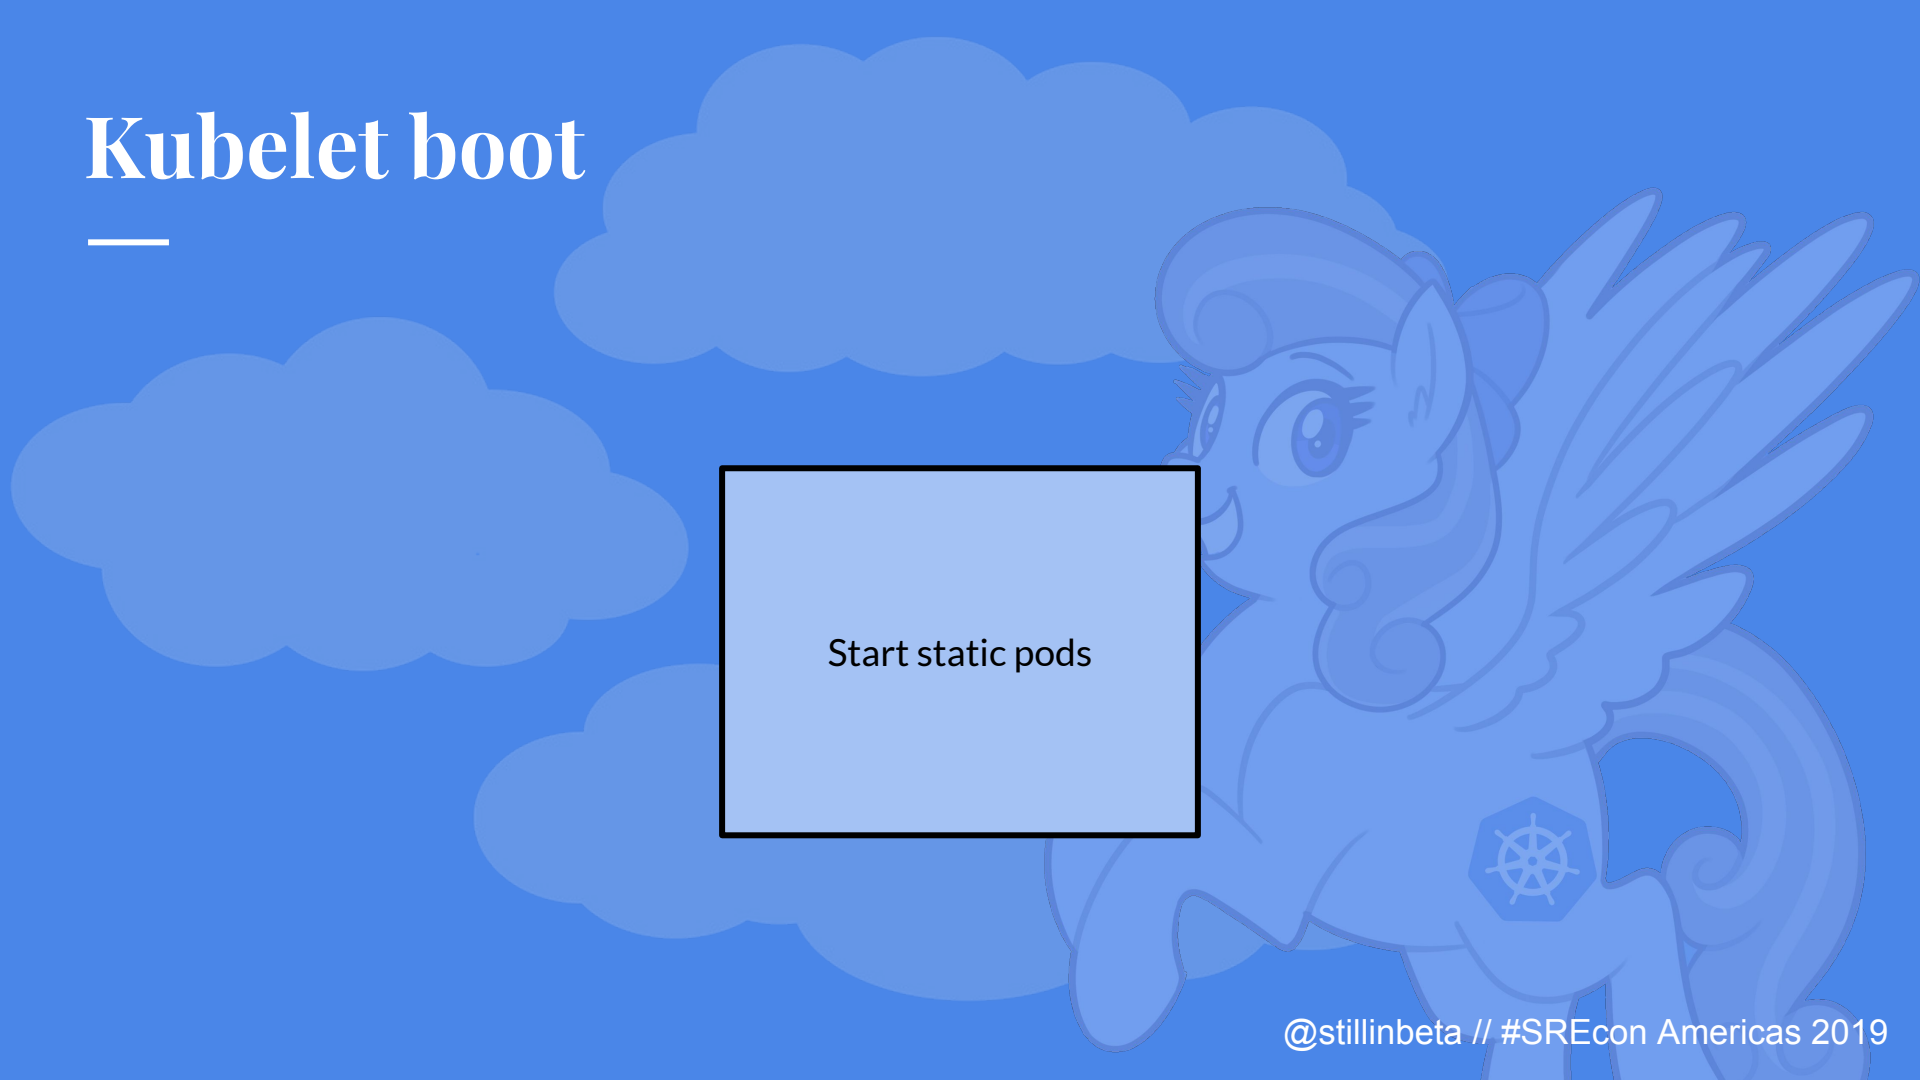

Tutorial 1!

http://bit.ly/srecon-cake1

kubeadm

kubeadm

Preflight Checks

- Avoid common problems
- Handy guard-rails
- Ignore at your own peril!

Tutorial 2!

http://bit.ly/srecon-cake2


etcd

From the Website

- Simple: well-defined, user-facing API (gRPC)
- Secure: automatic TLS with optional client cert authentication
- Fast: benchmarked 10,000 writes/sec
- Reliable: properly distributed using Raft

What you care about

- Key/value store
- *Optionally* distributed
- Only stateful part of kubernetes
- Secured by client certificates

Where do we put it?

- External?
- HA?
- Self-Hosted
- Single Node

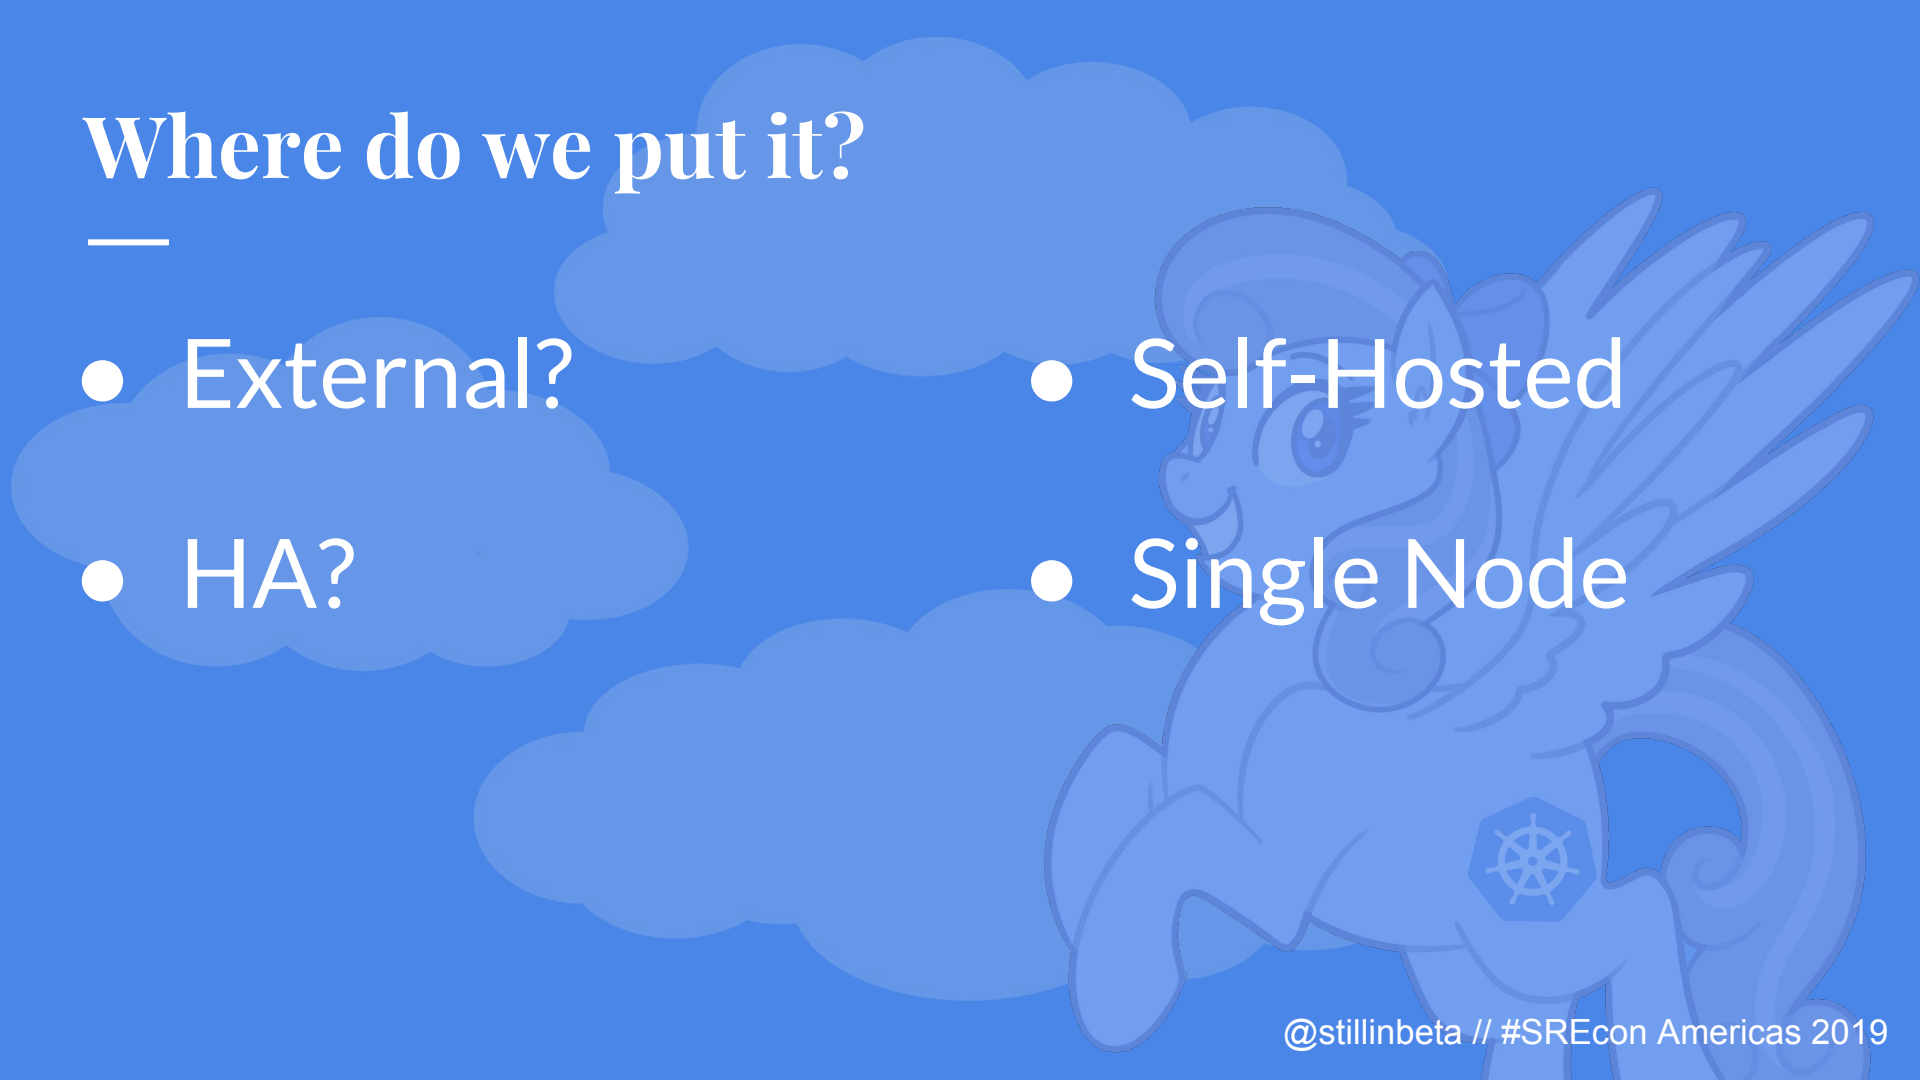

Setting up etcd

Tutorial 3!

http://bit.ly/srecon-cake3

apiserver

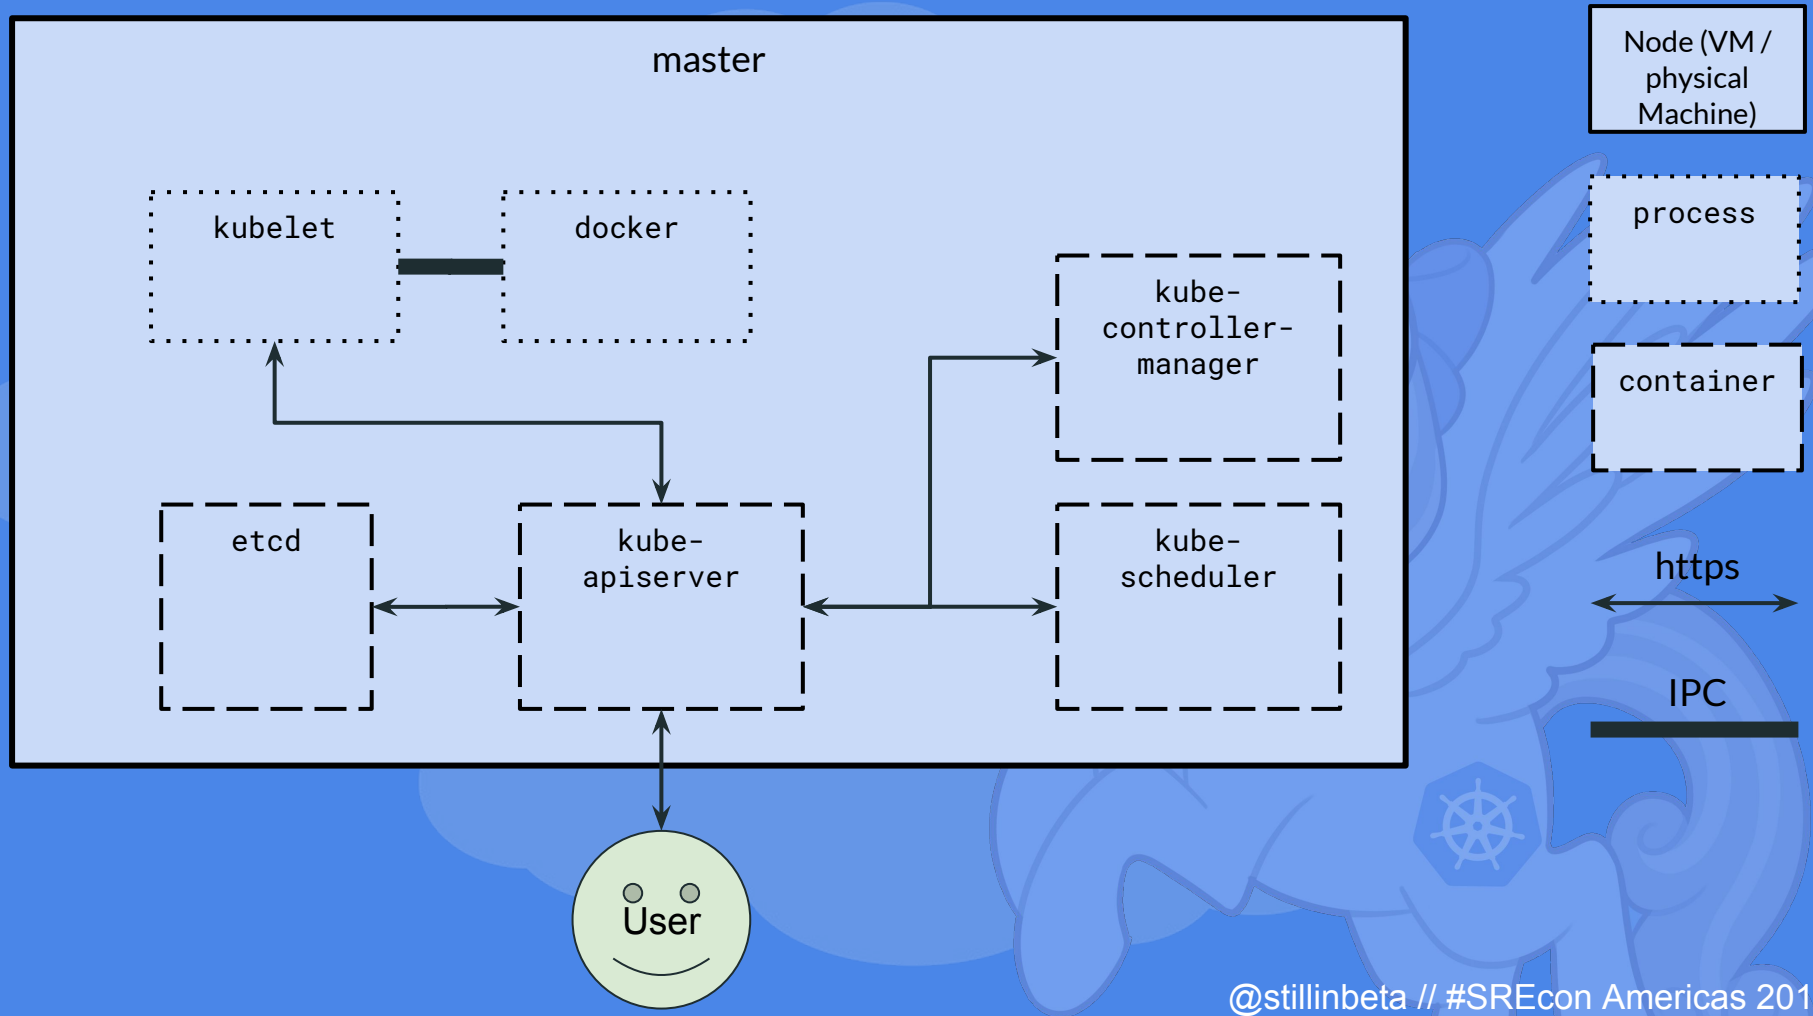
master
kubelet
docker
kube-controller-manager
etcd
kube-apiserver
kube-scheduler
User
Node (VM / physical Machine)
process
container
https
IPC

The centre of everything


kube-apiserver

kube-apiserver

- Exclusive control of etcd
- Handles all access control
- Stores some (but not all!) APIs
- Proxies everything

API Groups?


GET /apis/apps/v1/namespaces/default/deployments/myapp
API Group
API Version
Namespace
Resource
Resource Name

Custom Resource Definitions

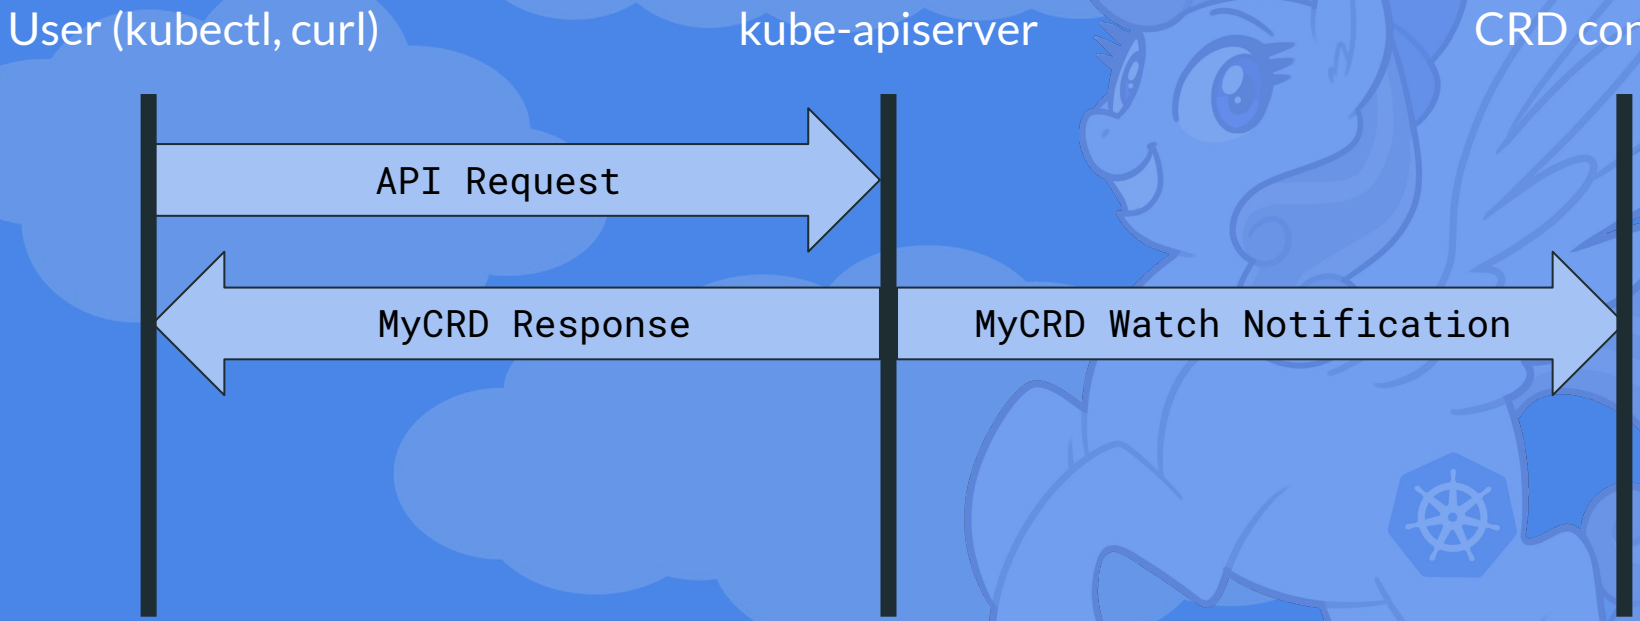

kubectl exec

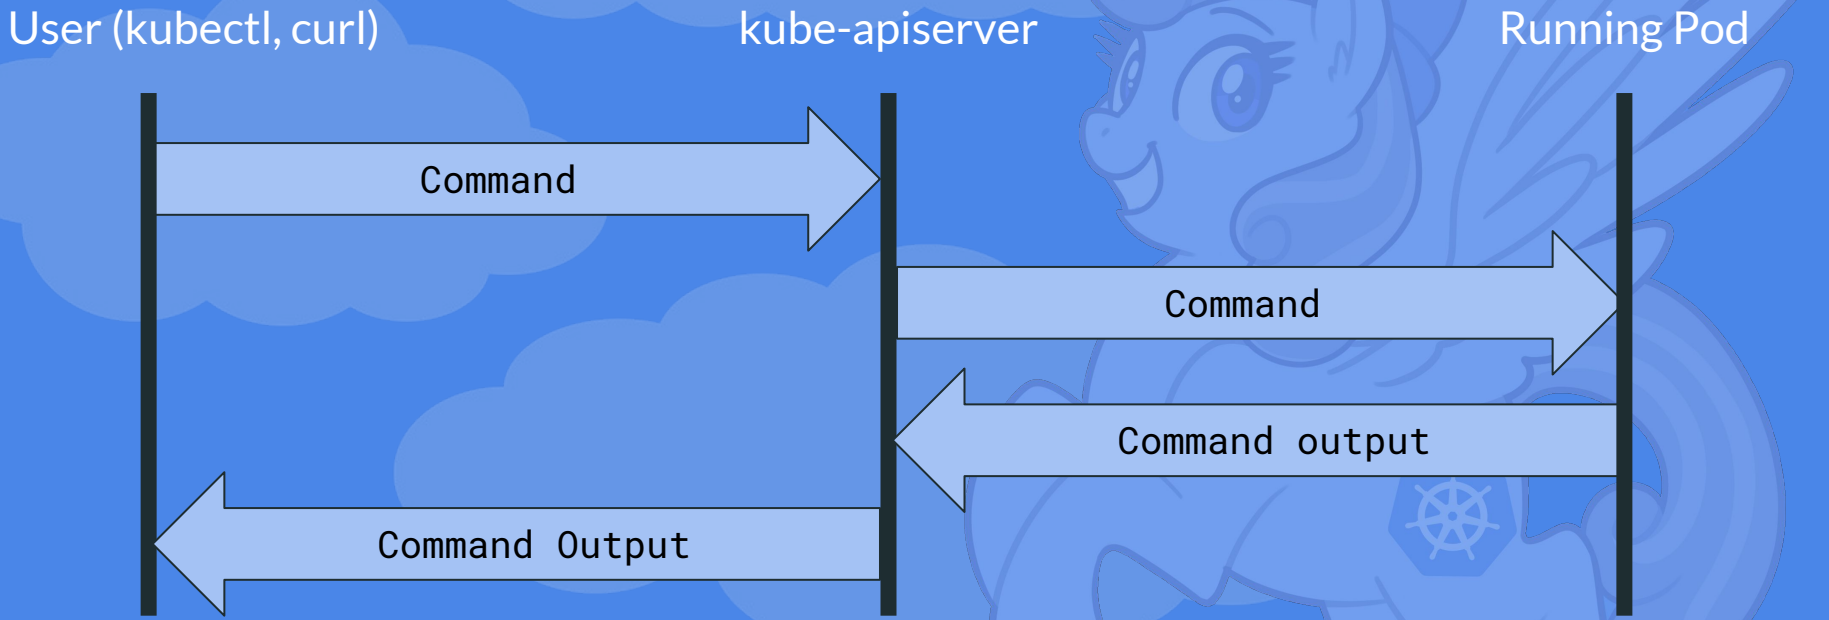

Tutorial 4!

http://bit.ly/srecon-cake4

kube-controller-manager

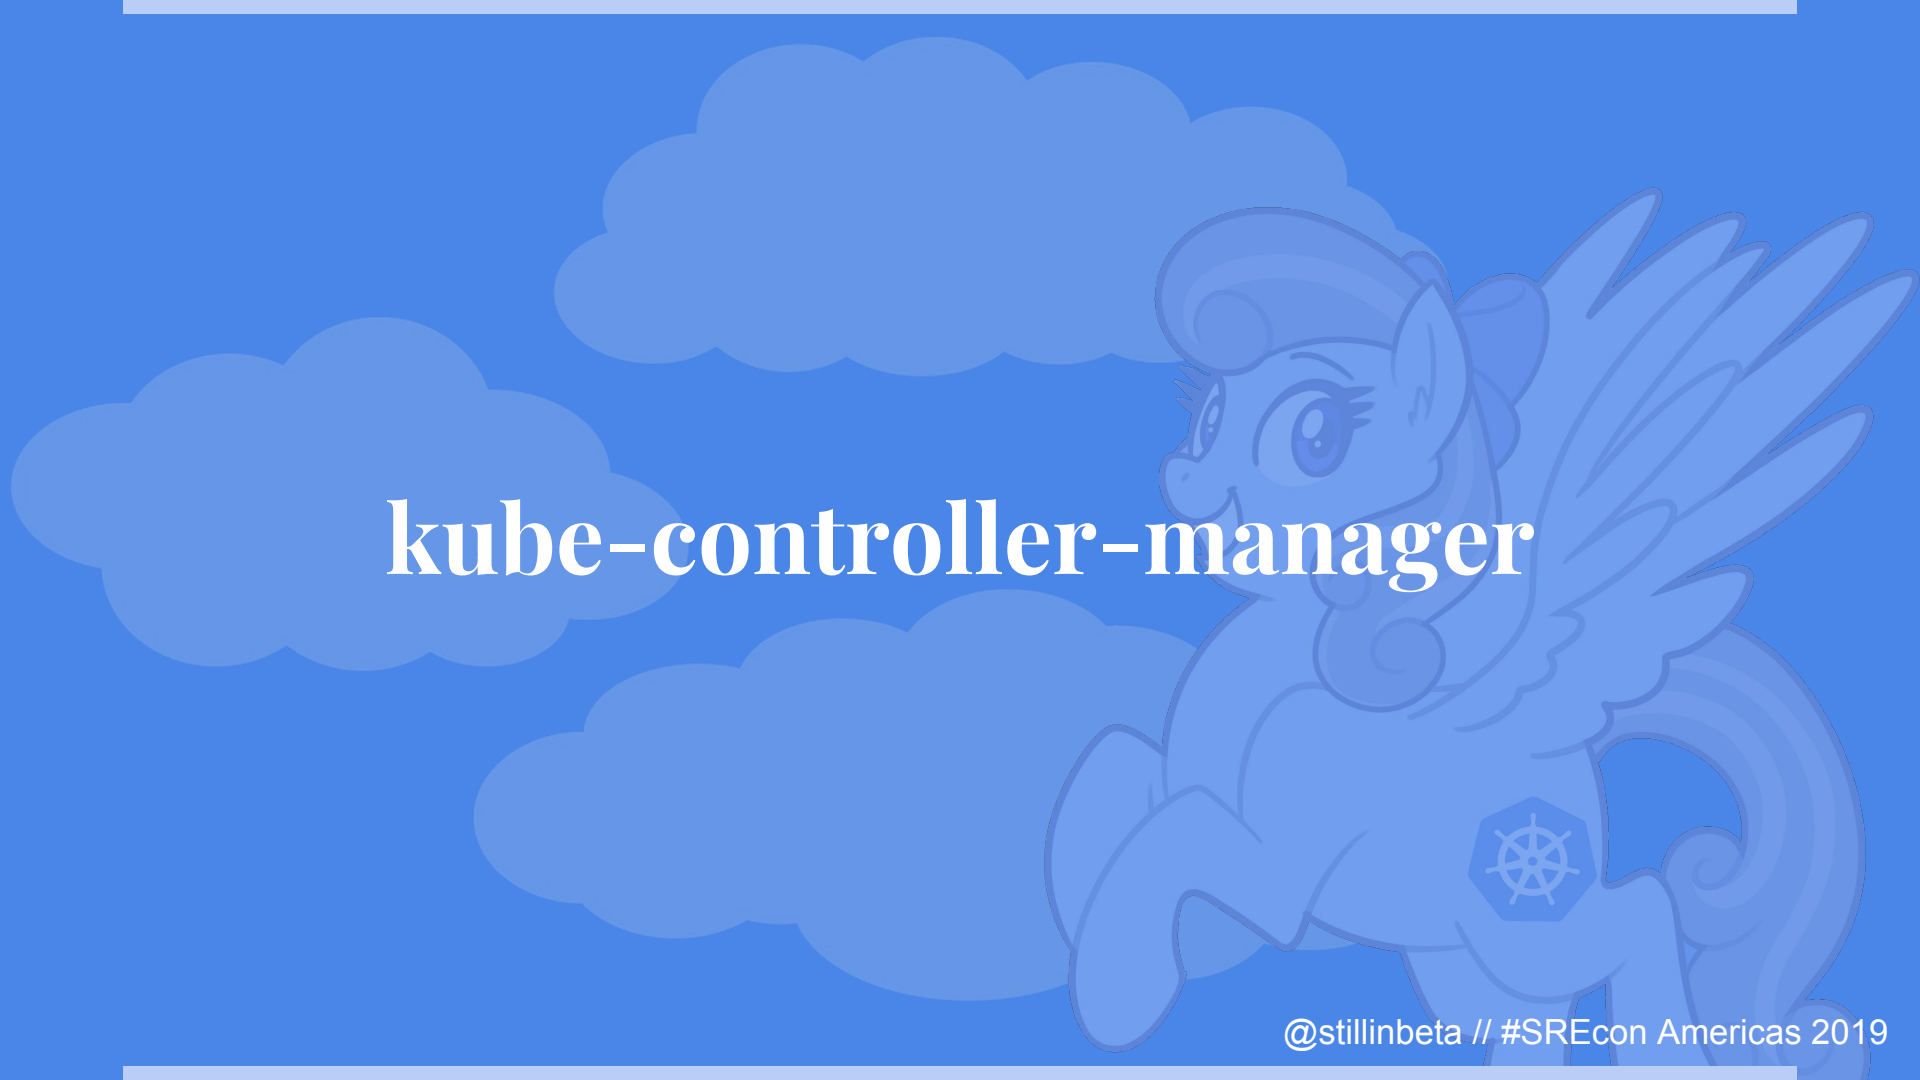

...that could mean anything

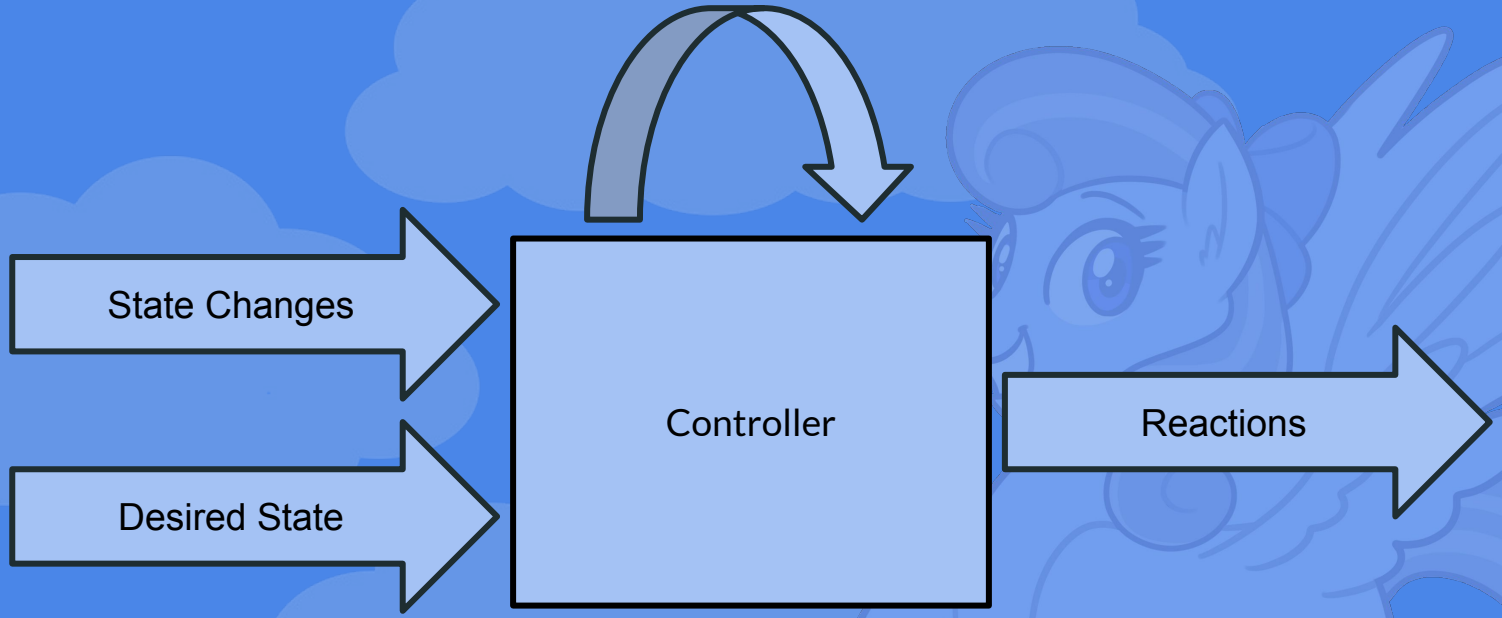
State Changes
Desired State
Controller
Reactions

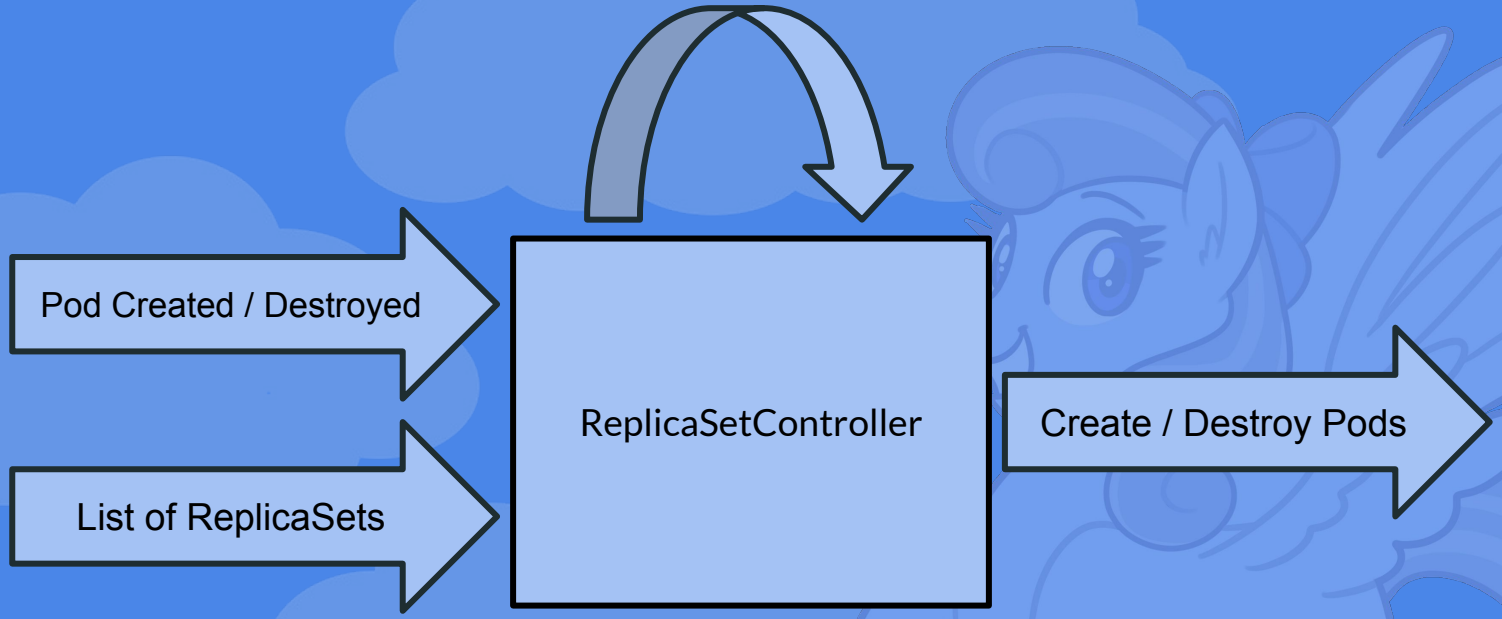
Pod Created / Destroyed
List of ReplicaSets
ReplicaSetController
Create / Destroy Pods

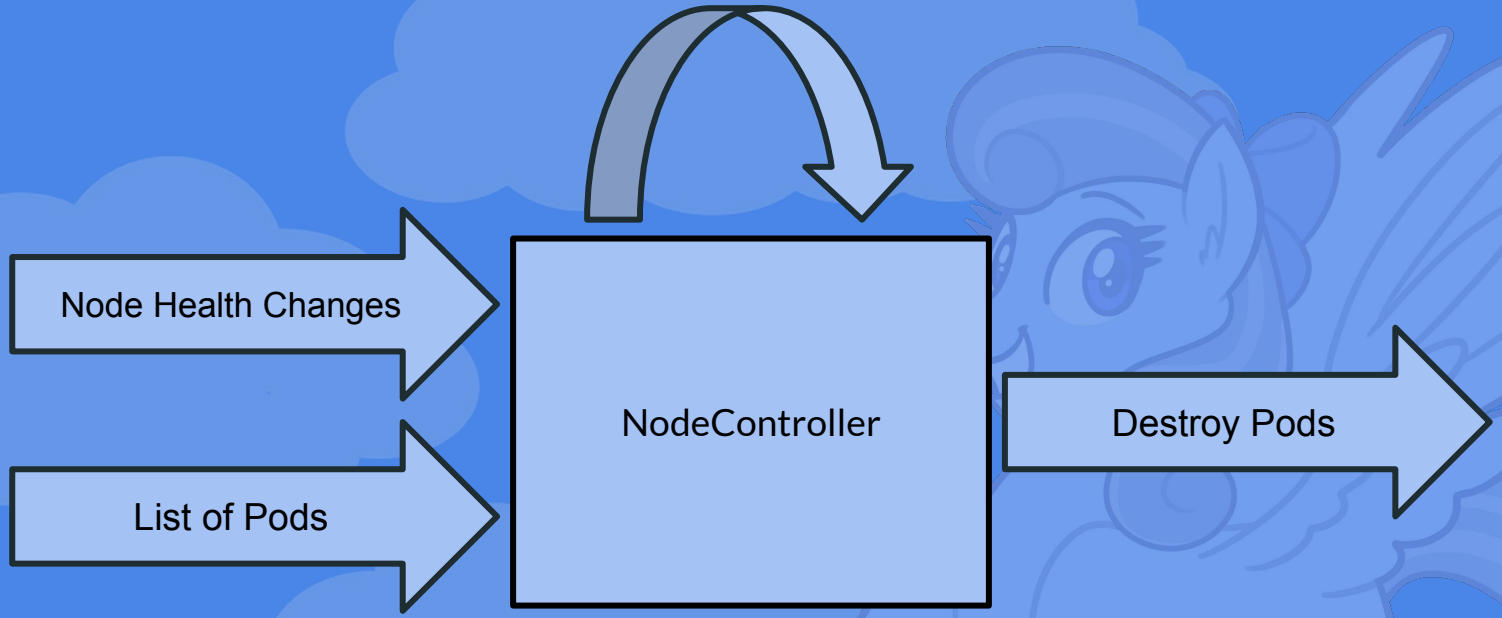
Node Health Changes
List of Pods
NodeController
Destroy Pods

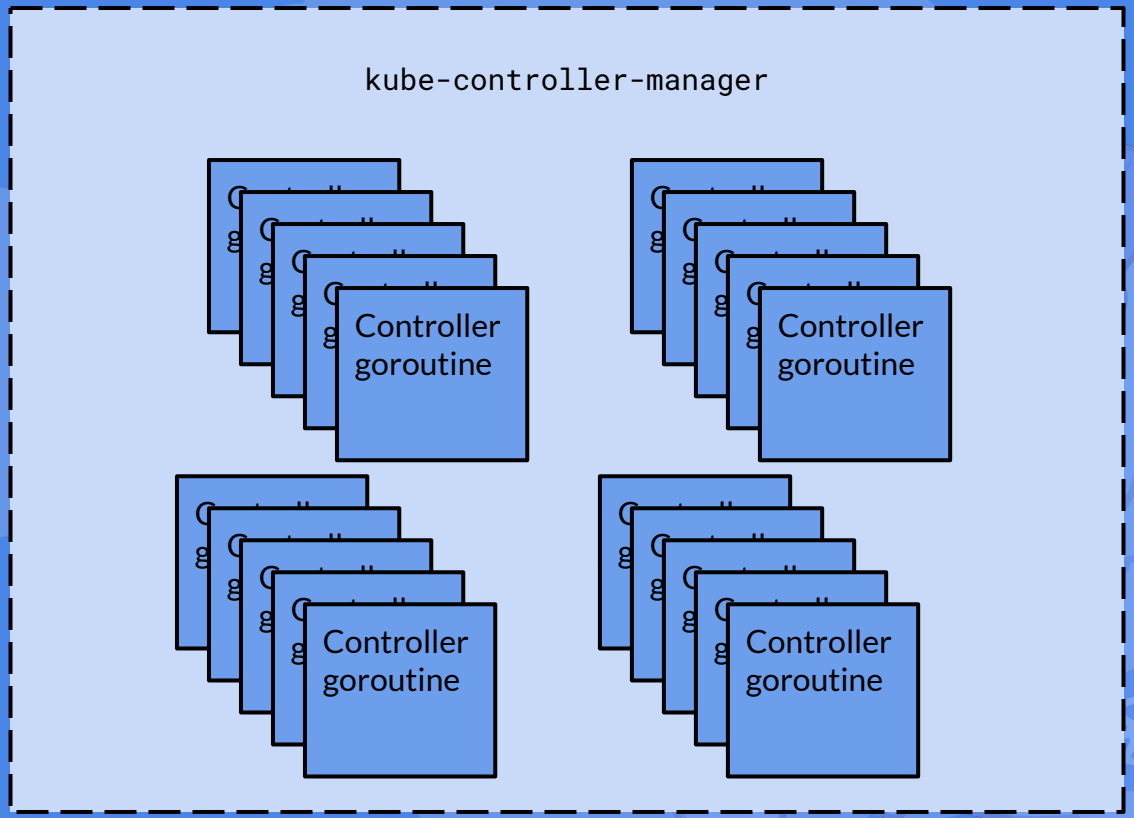
kube-controller-manager
Controller goroutine
Controller goroutine
Controller goroutine
Controller goroutine

Tutorial 5!

http://bit.ly/srecon-cake5

Scheduler

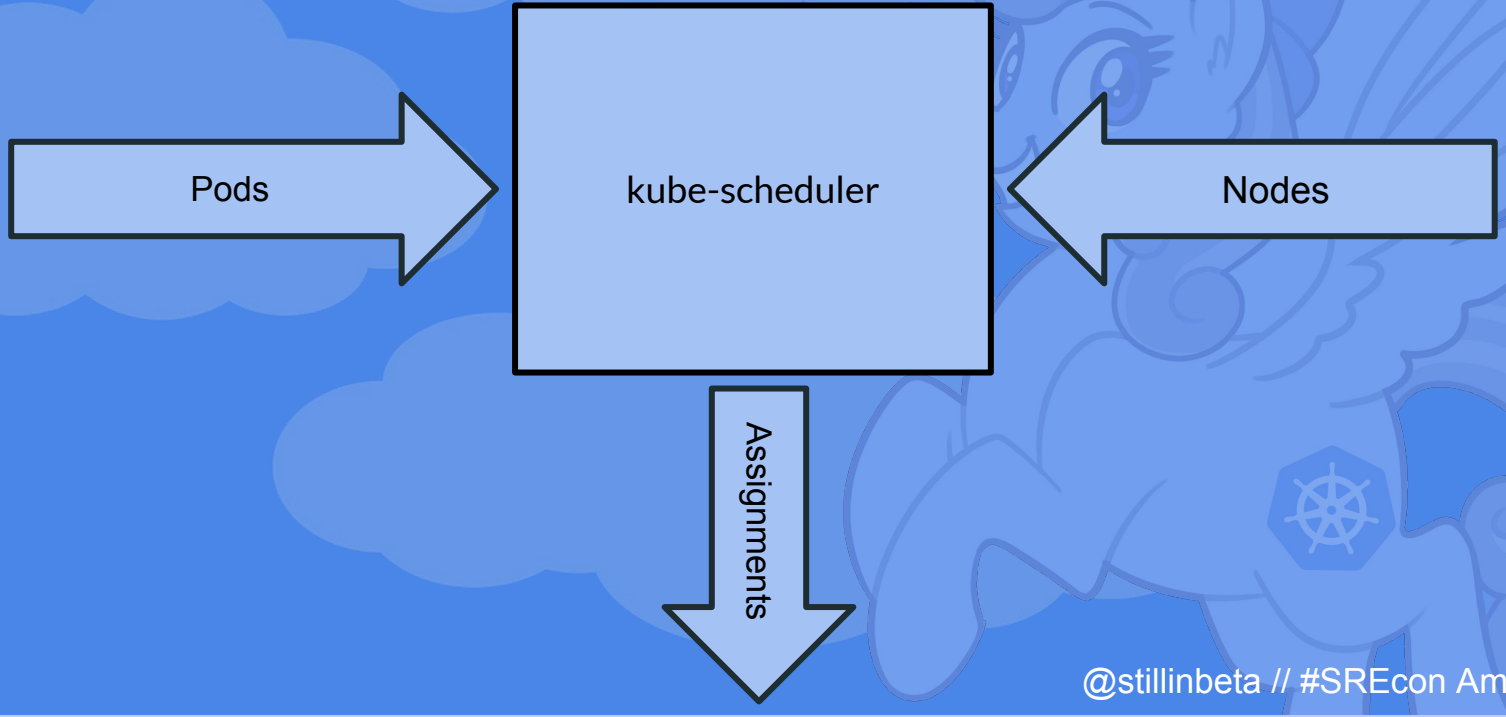
Pods
kube-scheduler
Nodes
Assignments

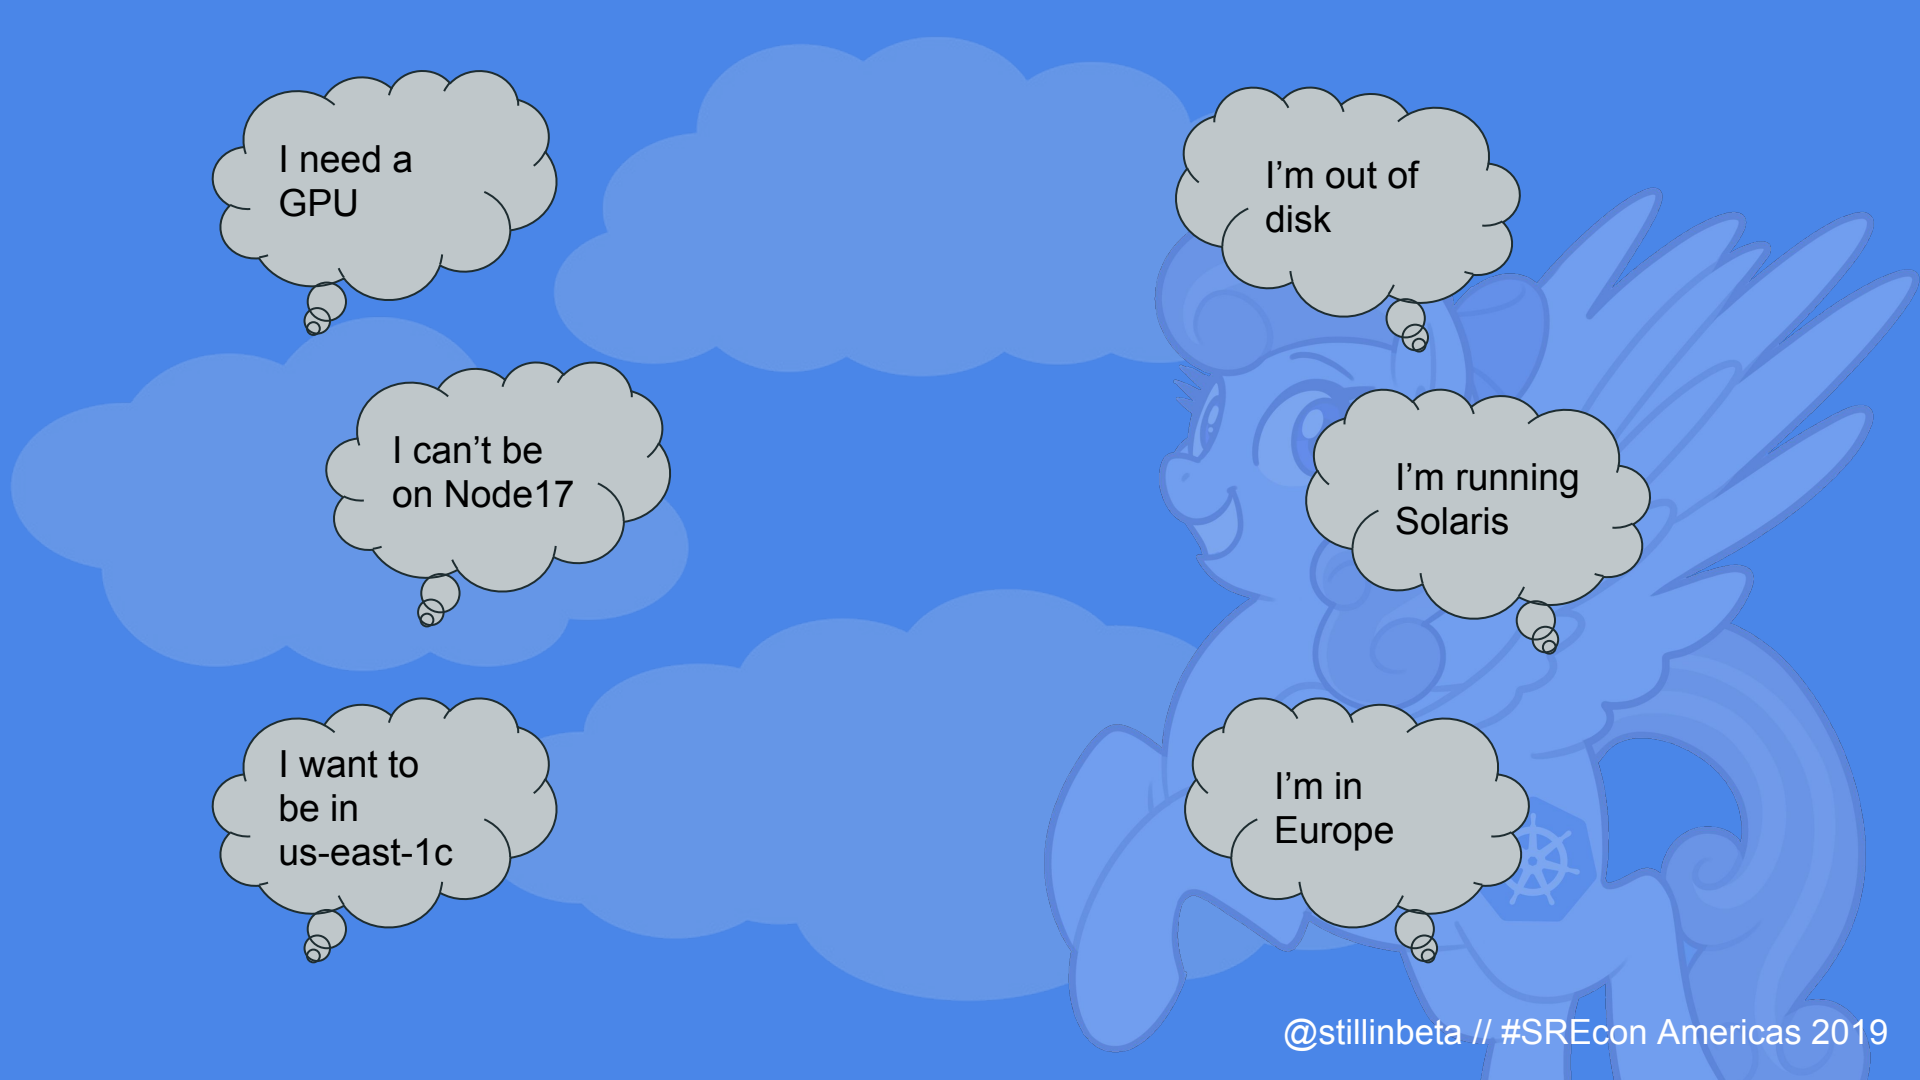
I need a GPU
I can't be on Node17
I want to be in us-east-1c
I'm out of disk
I'm running Solaris
I'm in Europe


I *prefer* Europe
I prefer not to run on the same node as pod17

Bin Packing

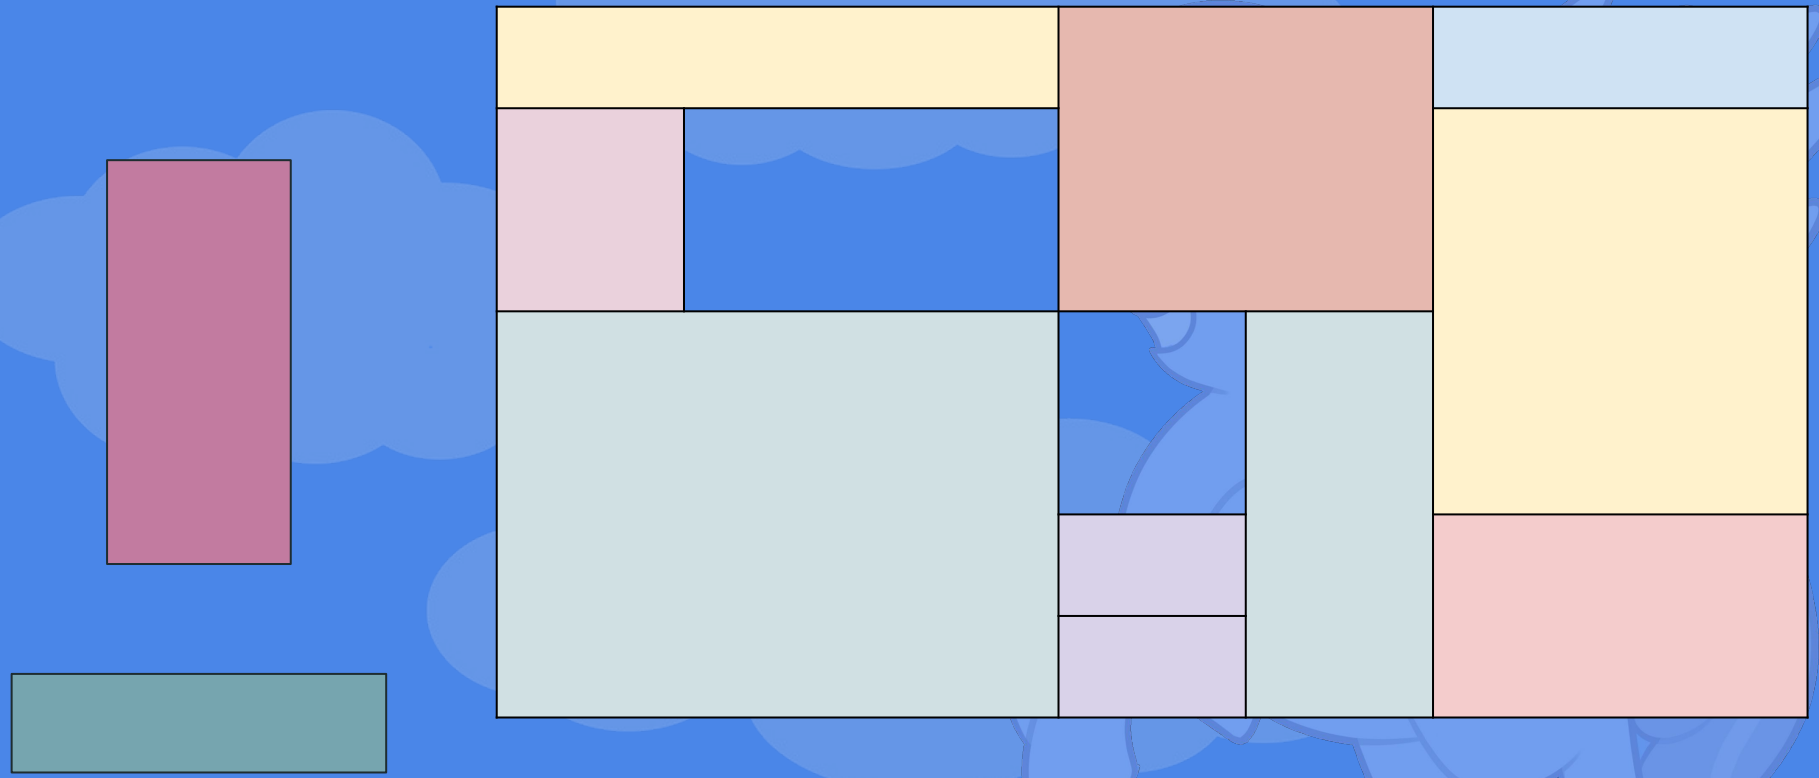

Tutorial 6!

http://bit.ly/srecon-cake6

What would be next?

- DNS
- Container Networking
- Proxying
- Multiple Nodes

Your Cluster is *Not:*

- Secure
- Production-Ready
- Resilient
- High-Availability
- Certified Kubernetes


But it is *yours.*

Questions? Concerns?

Come find me!
I'm the one with pink hair!

Twitter: @stillinbeta

slack.k8s.io: @liz

Github: github.com/stillinbeta